

H2 II. Representing and Manipulating Information

H3 §2.1 Information Storage

- **Bytes**: blocks of 8 bits, the smallest addressable unit of memory
- **Virtual memory**: machine-level program views memory as a very large array of bytes
- **Virtual address space**: set of all possible addresses
 - conceptual image presented to the machine-level program
 - actual implementation=DRAM + flash memory + disk storage + special hardware + OS

H5 ¶2.1.1 Hexadecimal Notation

- 1 byte = 8 bits
 - $00000000_2 \sim 11111111_2$
 - $0_{10} \sim 255_{10}$
 - Base-16 or **hexadecimal** is commonly used for convenience
 - $00_{16} \sim FF_{16}$
 - In C, numeric constants starting with **0x** or **0X** are interpreted as being in hexadecimal
 - characters 'A' ~ 'F' may be written in either upper or lowercase

H5 ¶2.1.2 Data Sizes

- **Word size**: nominal size of pointer data
 - virtual address is encoded by words
 - determines the **maximum size of the virtual address space**
 - machine with a w-bit word size gives the program access to at most 2^w bytes
- In recent years, widespread shift from 32-bit machines to **64-bit machines**
 - 32-bit: virtual address space $\approx 4\text{GB}$ (4×10^9 bytes)
 - 64-bit: virtual address space $\approx 16\text{EB}$ (1.84×10^{19} bytes)
 - Most 64-bit machines can also run programs compiled for use on 32-bit machines (**backward compatibility**)
 - 64-bit programs will only run on a 64-bit machine

```

1  linux> gcc -m32 prog.c // runs on either a 32-bit or a 64-
    bit machine
2
3  linux> gcc -m64 prog.c // runs only on a 64-bit machine

```

- C language supports multiple data formats for both integer and floating point data
 - exact numbers of bytes for some data types depends on how the program is compiled
 - most of the data types **encode signed values**, unless prefixed by the keyword **unsigned**
 - **char**: C standard does not guarantee it to be encoded as signed
 - use **signed char** to guarantee a 1-byte signed value
 - **pointer**: uses the full word size of the program

C Declaration		Bytes	
Signed	Unsigned	32-bit	64-bit
[signed] char	unsigned char	1	1
short	unsigned short	2	2
int	unsigned	4	4
long	unsigned long	4	8
int32_t	uint32_t	4	4
int64_t	uint64_t	8	8
char*		4	8
float		4	4
double		8	8

- ```

1 - For portability, make the program insensitive to the exact sizes
 of the different data types

```

## H5 ¶2.1.3 Addressing and Byte Ordering

- For program objects that span multiple bytes, need to establish two conventions:
  1. what the address of the object will be
  2. how we will order the bytes in memory
    - Two common conventions
      1. Little endian
        - store the object in memory ordered from least significant byte to most
        - Most Intel-compatible machines
      2. Big endian
        - store the object in memory ordered from most significant byte to least
        - IBM, Oracle (Sun Microsystems)
  - E.g. variable `x` of type `int` at address **0x100** has a hexadecimal value of **0x1234567**
    - Big endian

```
0x100 0x101 0x102 0x103
 01 23 45 67
```
    - Little endian

```
0x100 0x101 0x102 0x103
 67 45 23 01
```
  - recent microprocessor chips are **bi-endian** (configurable)
  - no technological reason to choose one byte ordering convention over the other
  - Byte ordering becomes an **issue** when...
    1. binary data are communicated over a **network** between different machines
      - make sure sending machine converts its internal representation to the network standard, while the receiving machine converts the network standard to its internal representation
    2. looking at the **byte sequences** representing integer data
      - inspecting machine-level programs

```
1 4004d3: 01 05 43 0b 20 00 add %eax, 0x200b45(%rip)
```

- line generated by a **disassembler**
  - : tool that determines the instruction sequence represented by an executable program file
  - adds a word of data to the value stored at an address computed by adding

**0x200b43** to the current value of the **program counter**, the address of the next instruction to be executed

3. programs are written that circumvent the normal type system
  - using **cast** or a **union** to allow an object to be referenced according to a different data type from which it was created

```
1 #include <stdio.h>
2
3 typedef unsigned char *byte_pointer;
4
5 void show_bytes(byte_pointer start, size_t len) {
6 int i;
7 for(i = 0; i < len; i++)
8 printf(" %.2x", start[i]);
9 printf("\n");
10 }
11
12 void show_int(int x) {
13 show_bytes((byte_pointer) &x, sizeof(int));
14 }
15
16 void show_float(float x) {
17 show_bytes((byte_pointer) &x, sizeof(float));
18 }
19
20 void show_pointer(void *x) {
21 show_bytes((byte_pointer) &x, sizeof(void *));
22 }
```

- in the code above, functions pass `show_bytes` a pointer `&x` to their argument `x`, casting the pointer to be of type **unsigned char \***
  - indicates to the compiler that the program should consider the pointer to be a sequence of bytes rather than to an object of the original data type

## H5 ¶2.1.5 Representing Code

- **instruction codings are different**
  - different machine types use different and incompatible instructions and encodings
  - even identical processors running different OS have differences in their coding conventions. --> not binary compatible
- Fundamental concept of computer system: **program**, from the perspective of the machine, **is simply a sequence of bytes**

## H5 ¶2.1.6 Introduction to Boolean Algebra

- **George Boole** observed that by encoding logic values **TRUE** and **FALSE** as binary values 1 and 0, he could formulate an algebra that captures the basic principles of logical reasoning
- Operations of Boolean algebra

|   |       |   |       |   |   |       |   |   |       |   |   |
|---|-------|---|-------|---|---|-------|---|---|-------|---|---|
| 1 | ~     |   | &     | 0 | 1 |       | 0 | 1 | ^     | 0 | 1 |
| 2 | ----- |   | ----- |   |   | ----- |   |   | ----- |   |   |
| 3 | 0     | 1 | 0     | 0 | 0 | 0     | 0 | 1 | 0     | 0 | 1 |
| 4 | 1     | 0 | 1     | 0 | 1 | 1     | 1 | 1 | 1     | 1 | 0 |

- ~: logical operation **NOT**
  - $\sim \text{TRUE} = \text{FALSE}$  /  $\sim \text{FALSE} = \text{TRUE}$
- &: logical operation **AND**
  - $p \& q == 1$  only when  $p = 1$  and  $q = 1$
- |: logical operation **OR**
  - $p | q == 1$  when either  $p = 1$  or  $q = 1$
- ^: logical operation **EXCLUSIVE-OR**
  - $p \wedge q == 1$  when either  $p = 1$  and  $q = 0$  or  $p = 0$  and  $q = 1$
- Boolean algebra still plays a central role in the design and analysis of digital systems
- Boolean algebra can be extended on **bit vectors**

|   |        |      |        |        |
|---|--------|------|--------|--------|
| 1 | 0110   | 0110 | 0110   |        |
| 2 | & 1100 | 1100 | ^ 1100 | ~ 1100 |
| 3 | ----   | ---- | ----   | ----   |
| 4 | 0100   | 1110 | 1010   | 0011   |

## H5 ¶2.1.7 Bit-Level Operations in C

- C supports **bitwise Boolean operations** applied to any "integral data type"

| C expression                 | Binary expression                          | Binary result            | Hexadecimal result |
|------------------------------|--------------------------------------------|--------------------------|--------------------|
| <code>~0x41</code>           | <code>~[0100 0001]</code>                  | <code>[1011 1110]</code> | <code>0xBE</code>  |
| <code>~0x00</code>           | <code>~[0000 0000]</code>                  | <code>[1111 1111]</code> | <code>0xFF</code>  |
| <code>0x69 &amp; 0x55</code> | <code>[0110 1001] &amp; [0101 0101]</code> | <code>[0100 0001]</code> | <code>0x41</code>  |
| <code>0x69   0x55</code>     | <code>[0110 1001]   [0101 0101]</code>     | <code>[0111 1101]</code> | <code>0x7D</code>  |

- commonly used to implement **masking** operations
  - mask**: bit pattern that indicates a selected set of bits within a word
    - e.g. **0xFF**: the lower-order byte of a word
- $\Rightarrow x \& 0xFF$  returns the least significant byte of  $x$  with all other bytes set to 0
  - if  $x = 0x89ABCDEF$ ,  $x \& 0xFF = 0x000000EF$

## H5 ¶2.1.8 Logical Operations in C

- C provides a set of **logical operators** `||`, `&&`, and `!`
  - correspond to the OR, AND, and NOT operations
- Logical operations treat any **nonzero argument** as **TRUE** and argument **0** as **FALSE**, then return either **1 (TRUE)** or **0 (FALSE)**

## H5 ¶2.1.9 Shift Operations in C

- C provides a set of **shift** operations for **shifting bit patterns** to the left and to the right
  - if  $x = [x_{w-1}, x_{w-2}, \dots, x_0]$ ,  $x \ll k$  yields  $[x_{w-1-k}, x_{w-k-2}, \dots, x_0, 0, \dots, 0]$ 
    - $x$  is shifted **k** bits to the left, dropping off the **k** most significant bits and filling the right end with **k** zeros
  - Shift operations associate from left to right
    - $x \ll j \ll k$  is equivalent to  $(x \ll j) \ll k$
- Note that machines support **two forms of right shift**  $[x \gg k]$ 
  - Logical**: fills the left end with **k** zeros
    - $[0, \dots, 0, x_{w-1}, x_{w-2}, \dots, x_k]$
  - Arithmetic**: fills the left end with **k** repetitions of the most significant bit
    - $[x_{w-1}, \dots, x_{w-1}, x_{w-1}, x_{w-2}, \dots, x_k]$
    - useful for operating on signed integer data

- C standards do not precisely define which type of right shift should be used with signed numbers -- either arithmetic or logical shifts may be used
  - Almost all use **arithmetic right shifts for signed data**
  - **logical right shifts for unsigned data**

### H3 §2.2 Integer Representations

- two different ways bits can be used to encode integers
  1. only representing nonnegative numbers
  2. representing negative, zero, and positive numbers
- strongly related both in their mathematical properties and their machine-level implementations

#### H5 ¶2.2.1 Integral Data Types

- C supports a variety of **integral** data types -- ones that represent finite ranges of **integers**
  - each type can specify a size with keywords, **char, short, long**
  - can indicate whether the represented numbers are all **nonnegative** (declared as **unsigned**) or **negative** (by **default**)
- different sizes allow different ranges of values to be represented
- **long** is the only **machine-dependent range** indicator
- **Note that** ranges are **not symmetric**
  - range of negative numbers extends one further than the range of positive numbers

#### H5 ¶2.2.2 Unsigned Encodings

- consider an integer data type of **w** bits and write a bit vector as  $\vec{x}$

For vector  $\vec{x} = [x_{w-1}, x_{w-2}, \dots, x_0]$ :

$$B2U_w(\vec{x}) := \sum_{i=0}^{w-1} x_i 2^i$$

where  $B2U_w$  is a function that interprets **binary to unsigned**

for example,

$$B2U_4([0001]) = 0 \cdot 2^3 + 0 \cdot 2^2 + 0 \cdot 2^1 + 1 \cdot 2^0 = 0 + 0 + 0 + 1 = 1$$

$$B2U_4([0101]) = 0 \cdot 2^3 + 1 \cdot 2^2 + 0 \cdot 2^1 + 1 \cdot 2^0 = 0 + 4 + 0 + 1 = 5$$

$$B2U_4([1011]) = 1 \cdot 2^3 + 0 \cdot 2^2 + 1 \cdot 2^1 + 1 \cdot 2^0 = 8 + 0 + 2 + 1 = 11$$

$$B2U_4([1111]) = 1 \cdot 2^3 + 1 \cdot 2^2 + 1 \cdot 2^1 + 1 \cdot 2^0 = 8 + 4 + 2 + 1 = 15$$

- Function  $B2U_w$  is a bijection
- $UMax_w := \sum_{i=0}^{w-1} 2^i = 2^w - 1$

## H5 ¶2.2.3 Two's-Complement Encodings

- to represent negative values, **Two's complement form** uses the **most significant bit** of the word to have **negative weight**

For vector  $\vec{x} = [x_{w-1}, x_{w-2}, \dots, x_0]$ :

$$B2T_w(\vec{x}) := -x_{w-1}2^{w-1} + \sum_{i=0}^{w-2} x_i 2^i$$

where  $B2T_w$  is a function that interprets **binary to two's complement**

for example,

$$B2T_4([0001]) = -0 \cdot 2^3 + 0 \cdot 2^2 + 0 \cdot 2^1 + 1 \cdot 2^0 = 0 + 0 + 0 + 1 = 1)$$

$$B2T_4([0101]) = -0 \cdot 2^3 + 1 \cdot 2^2 + 0 \cdot 2^1 + 1 \cdot 2^0 = 0 + 4 + 0 + 1 = 5)$$

$$B2T_4([1011]) = -1 \cdot 2^3 + 0 \cdot 2^2 + 1 \cdot 2^1 + 1 \cdot 2^0 = -8 + 0 + 2 + 1 = -5)$$

$$B2T_4([1111]) = -1 \cdot 2^3 + 1 \cdot 2^2 + 1 \cdot 2^1 + 1 \cdot 2^0 = -8 + 4 + 2 + 1 = -1)$$

- Function  $B2T_w$  is a bijection
- $TMin_w := -2^{w-1}$
- $TMax_w := \sum_{i=0}^{w-2} 2^i = 2^{w-1} - 1$

- Note that...

1. **Two's complement range is asymmetric:**  $|TMin| = |TMax| + 1$

- half the bit patterns represent negative numbers, while half represent nonnegative numbers (including 0)

2. **Maximum unsigned value is just over twice the maximum two's complement value:**  $UMax = 2TMax + 1$

3. **-1 has the same bit representation as  $UMax$**



## H5 ¶2.2.4 Conversions between Signed and Unsigned

- C allows casting between different numeric data types
  - suppose **x** is declared as **int** and **u** as **unsigned**
    - **(unsigned) x** : converts x to an unsigned value
    - **(int) u**: converts u to signed integer
- Conversions are done based on a **bit-level perspective**
  - with the same word size, the **numeric values might change** , but the **bit patterns stay the same**

given an integer  $x$  in the range  $0 \leq x < UMax_w$ , the function  $U2B_w(x)$  gives the unique  $w$ -bit unsigned representation of  $x$ .

Similarly, when  $x$  is in the range  $TMin_w \leq x \leq TMax_w$ , the function  $T2B_w(x)$  gives the unique  $w$ -bit two's-complement representation of  $x$ .

Now, define the function  $T2U_w(x) := B2U_w(T2B_w(x))$ .

this function takes a number between  $TMin_w$  and  $TMax_w$  and yields a number between 0 and  $UMax_w$

the two numbers have identical bit representations, except that the argument has a two's-complement representation while the result is unsigned.

Similar for the function  $U2T_w(x) := B2T_w(U2B_w(x))$

Hence, **conversion from two's-complement to unsigned**

For  $x$  such that  $TMin_w \leq x \leq TMax_w$  :

$$T2U_w(x) = \begin{cases} x + 2^w, & x < 0 \\ x, & x \geq 0 \end{cases}$$

For example,  $T2U_{16}(-12,345) = -12,345 + 2^{16} = 53,191$

$$T2U_w(-1) = -1 + 2^w = UMax_w$$

Similarly, **conversion from unsigned to two's-complement**

For  $u$  such that  $0 \leq u \leq UMax_w$ :

$$U2T_w(u) = \begin{cases} u, & u \leq TMax_w \\ u - 2^w, & u > TMax_w \end{cases}$$

## H5 ¶2.2.5 Signed vs. Unsigned in C

- almost all machines use two's-complement
  - most numbers are **signed by default**
  - needs to add 'U' or 'u' as a suffix to create unsigned constants
    - 12345U or 0x1A2Bu
- **explicit casting**

```
1 int tx, ty;
2 unsigned ux, uy;
3
4 tx = (int) ux;
5 uy = (unsigned) ty;
```

- **implicit casting**

```
1 int tx, ty;
2 unsigned ux, uy;
3
4 tx = ux; // cast to signed
5 uy = ty; // cast to unsigned
```

- **using printf**

```
1 int x = -1;
2 unsigned u = 2147483648; // 2^31
3
4 printf("x = %u = %d\n", x, x);
5 printf("u = %u = %d\n", u, u);
6
7 /* On 32-bit machine, it will print
8 x = 4294967295 = -1
9 y = 2147483648 = -2147483648 */
```

- When an operation is performed where **one operand is signed and the other is unsigned**,
  - **C implicitly casts the signed argument to unsigned** and performs the operations assuming the numbers are **nonnegative**
  - quite accurate for standard arithmetic operations, but..

- **weird results** for relational operators  $<$  and  $>$ 
  - e.g.  $-1 < 0U$  returns **False**
    - because C casts  $-1$  to  $4294967295U$

## H5 ¶2.2.6 Expanding the Bit Representation of a Number

- Conversion between integers **having different word sizes** while retaining **the same numeric value**
  - **may not be possible** when the **destination data type is too small** to represent the desired value
  - **smaller to larger data type** should **always be possible**
- 1. **Zero extension**: for converting an unsigned number to a larger data type
  - **add leading zeros**

Define bit vectors  $\vec{u} = [u_{w-1}, u_{w-2}, \dots, u_0]$  of width  $w$  and  $\vec{u}' = [0, \dots, 0, u_{w-1}, u_{w-2}, \dots, u_0]$  of width  $w'$ , where  $w' > w$ . Then  $B2T_w(\vec{x}) = B2T_{w'}(\vec{x}')$ .

2. **Sign extension**: for converting a two's-complement number to a larger data type
  - **add copies of the most significant bit**

Define bit vectors  $\vec{x} = [x_{w-1}, x_{w-2}, \dots, x_0]$  of width  $w$  and  $\vec{x}' = [x_{w-1}, x_{w-1}, \dots, x_{w-1}, x_{w-2}, \dots, x_0]$  of width  $w'$ , where  $w' > w$ . Then  $B2T_w(\vec{x}) = B2T_{w'}(\vec{x}')$ .

```
1 // When run as a 32-bit program on a big-endian machine that
 uses a two's complement representation,
2
3 sx = -12345: cf c7
4 usx = 53191: cf c7
5 x = -12345: ff ff cf c7 // ff ff = 1111..1111
6 ux = 53191: 00 00 cf c7
```

## H5 ¶2.2.7 Truncating Numbers

- Used to reduce the number of bits representing a number

```

1 int x = 53191;
2 short sx = (short) x; // -12345
3 int y = sx; // -12345

```

- casting x to be short will truncate a 32-bit int to a 16-bit short
- When truncating a  $w$ -bit number  $\vec{x} = [x_{w-1}, x_{w-2}, \dots, x_0]$  to a  $k$ -bit number, **drop the high-order  $w - k$  bits**
  - can alter its value -- **OVERFLOW!**

- **Truncation of an unsigned number**

Let  $\vec{x}$  be the bit vector  $[x_{w-1}, x_{w-2}, \dots, x_0]$ , and let  $\vec{x}'$  be the result of truncating it to  $k$  bits:  $\vec{x}' = [x_{k-1}, x_{k-2}, \dots, x_0]$ .

Let  $x = B2U_w(\vec{x})$  and  $x' = B2U_k(\vec{x}')$ . Then  $x' = x \bmod 2^k$

- all of the bits that were truncated have weights of the form  $2^i$ , where  $i \geq k$

- **Truncation of a two's-complement number**

Let  $\vec{x}$  be the bit vector  $[x_{w-1}, x_{w-2}, \dots, x_0]$ , and let  $\vec{x}'$  be the result of truncating it to  $k$  bits:  $\vec{x}' = [x_{k-1}, x_{k-2}, \dots, x_0]$ .

Let  $x = B2T_w(\vec{x})$  and  $x' = B2T_k(\vec{x}')$ . Then  $x' = U2T_k(x \bmod 2^k)$

- Applying function  $U2T_k$  will have the effect of converting the most significant bit  $x_{k-1}$  from having weight  $2^{k-1}$  to having weight  $-2^{k-1}$
- For example,

Converting  $x = 53,191$  from *int* to *short*.

Since  $2^{16} = 65,536 \geq x$ , we have  $x \bmod 2^{16} = x$ .

But since we need to convert it to a 16-bit two's complement number, we get

$x' = 53,191 - 65,536 = -12,345$

## H5 ¶2.2.8 Advice on Signed versus Unsigned

- Implicit casting of signed to unsigned leads to some non-intuitive behavior
  - program bugs & difficult to identify it
- To avoid errors or vulnerabilities...
  1. **NEVER use unsigned numbers**
    - few languages other than C support unsigned integers
    - other language designers viewed unsigned integers as more trouble than

they are worth

2. or use it for **collections of bits with no numeric interpretation**
  - **flags** describing various Boolean conditions
  - implementing mathematical packages for modular arithmetic & multiprecision arithmetic

### H3 §2.3 Integer Arithmetic

- adding two positive numbers can yield a negative result
- $x-y$  can yield something other than  $x-y < 0$
- **NEED to understand Computer arithmetic** to write more reliable codes

### H5 ¶2.3.1 Unsigned Addition

Consider two nonnegative integers  $x$  and  $y$ , such that  $0 \leq x, y < 2^w$ .

Each of these values can be represented by a  $w$ -bit unsigned number.

However, if we compute their sum,  $0 \leq x + y \leq 2^{w+1} - 2$ .

**Representing this sum could require  $w + 1$  bits**

- **"Word size inflation"** : cannot place any bound on the word size required to fully represent the results of arithmetic operations

Now, let  $+_w^u$  for arguments  $x$  and  $y$ , where  $0 \leq x, y < 2^w$ , be the result of truncating the integer sum  $x + y$  to be  $w$  bits long and viewing the result as an unsigned number

- form of modular arithmetic: computing the sum modulo  $2^w$  by discarding any bits with weight greater than  $2^{w-1}$

Then, **Unsigned Addition** can be formularized as...

For  $x$  and  $y$  such that  $0 \leq x, y < 2^w$ :

$$x +_w^u y = \begin{cases} x + y, & x + y < 2^w \\ x + y - 2^w, & 2^w \leq x + y < 2^{w+1} \end{cases} \quad \begin{matrix} \text{Normal} \\ \text{Over flow} \end{matrix}$$

In addition, to **detect overflow** of unsigned additions,

For  $x$  and  $y$  in the range  $0 \leq x, y \leq UMax_w$ , let  $s := x +_w^u y$ .

Then the computation of  $s$  overflowed if and only if  $s < x$  (or equivalently,  $s < y$ )

E.g.  $9 +_4^u 12 = 5$ .  $\implies$  **OVERFLOW!** ( $\because 5 < 9$ )

$(1001_2 + 1100_2 = 10101_2 \Rightarrow 0101_2)$  ( $\because$  word size = 4)

Similarly, for **Unsigned negation**,

For any number  $x$  such that  $0 \leq x < 2^w$ , its  $w$ -bit unsigned negation  $-_w^u x$  is given by the following:

$$-_w^u x = \begin{cases} x, & x = 0 \\ 2^w - x, & x > 0 \end{cases}$$

E.g.  $-_4^u 4 = 12$

(  
 $-0100_2 = 1100_2 = (-1) \cdot 1 \cdot 2^3 + 1 \cdot 2^2 + 0 \cdot 2^1 + 0 \cdot 2^0 = -8 + 4 + 0 + 0 = -4$ )

But since we are taking unsigned negation,  $1100_2 = 8 + 4 + 0 + 0 = 12$

## H5 ¶2.3.2 Two's-Complement Addition

- for two's complement addition, results can be either **too large** (positive) or **too small** (negative) to represent

Let us define  $x +_w^t y$  be the result of **truncating** the integer sum  $x + y$  to be  $w$  bits long.

For integer values  $x$  and  $y$  in the range  $-2^{w-1} \leq x, y \leq 2^{w-1} - 1$ :

$$x +_w^t y = \begin{cases} x + y - 2^w, & 2^{w-1} \leq x + y & \text{Positive overflow} \\ x + y, & -2^{w-1} \leq x + y < 2^{w-1} & \text{Normal} \\ x + y + 2^w, & x + y < -2^{w-1} & \text{Negative overflow} \end{cases}$$

- when  $x + y$  exceeds  $TMax_w$ ,  $\implies$  **positive overflow!**

e.g.  $x = 5 = [0101], y = [0101] \Rightarrow x + y = 10 = [01010]$ ,  $x +_4^t y = -6 = [1010]$

$$10 - 2^4 = -6$$

- when  $x - y$  is less than  $TMin_w \Rightarrow$  **negative overflow!**

e.g.  $x = -8 = [1000]$ ,  $y = -5 = [1011] \Rightarrow x + y = -13 = [10011]$ ,  $x +_w^t y = 3 = [0011]$   
 $-13 + 2^4 = 3$

To **detect overflow** in two's-complement addition

For  $x$  and  $y$  in the range  $TMin_w \leq x, y \leq TMax_w$ , let  $s := x +_w^t y$ .

Then the computation of  $s$  has had positive overflow if and only if  $x > 0$  and  $y > 0$  but  $s \leq 0$ .

The computation has had negative overflow if and only if  $x < 0$  and  $y < 0$  but  $s \geq 0$ .

## H5 ¶2.3.3 Two's-Complement Negation

- Every number  $x$  in the range  $TMin_w \leq x \leq TMax_w$  has an additive inverse under  $+_w^t$ , which we denote  $-_w^t x$  as follows:

For  $x$  in the range  $TMin_w \leq x \leq TMax_w$ , its two's-complement negation  $-_w^t x$  is given by the formula

$$-_w^t x = \begin{cases} TMin_w, & x = TMin_w \\ -x, & x > TMin_w \end{cases}$$

Note that  $TMin_w + TMin_w = -2^{w-1} + -2^{w-1} = -2^w \Rightarrow$  **Negative overflow!**

- Bit-level representation can be used to find two's-complement negation examples with a 4-bit word size:

| $\vec{x}$   | $\sim \vec{x}$ | $incr(\sim \vec{x})$ |
|-------------|----------------|----------------------|
| $[0101]=5$  | $[1010]=-6$    | $[1011]=-5$          |
| $[0111]=7$  | $[1000]=-8$    | $[1001]=-7$          |
| $[0000]=0$  | $[1111]=-1$    | $[0000]=0$           |
| $[1000]=-8$ | $[0111]=7$     | $[1000]=-8$          |

## H5 ¶2.3.4 Unsigned Multiplication

Integers  $x$  and  $y$  in the range  $0 \leq x, y \leq 2^w - 1$  can be represented as  $w$ -bit unsigned numbers, but their product  $x \cdot y$  can range between 0 and  $(2^w - 1)^2 = 2^{2w} - 2^{w+1} + 1$ , requiring  $2w$  bits to represent

Let's define  $x *^u_w y$  be the  $w$ -bit value given by the low-order  $w$  bits of the  $2w$ -bit integer product.

Then, **Unsigned multiplication** can be formularized as:

For  $x$  and  $y$  such that  $0 \leq x, y \leq UMax_w$ :

$$x *^u_w y = (x \cdot y) \bmod 2^w$$

## H5 ¶2.3.5 Two's-Complement Multiplication

Integers  $x$  and  $y$  in the range  $-2^{w-1} \leq x, y \leq 2^{w-1} - 1$  can be represented as  $w$ -bit two's complement numbers, but their product  $x \cdot y$  can range between  $-2^{w-1} \cdot (2^{w-1} - 1) = -2^{2w-2} + 2^{w-1}$  and  $-2^{w-1} \cdot -2^{w-1} = 2^{2w-2}$ , requiring as many as  $2w$  bits to represent in two's complement form

Let's define  $x *^t_w y$  be the  $w$  bit result after truncating the  $2w$ -bit product.

Then, **Two's-complement multiplication** can be formularized as:

For  $x$  and  $y$  such that  $TMin_w \leq x, y \leq TMax_w$ :

$$x *^t_w y = U2T_w((x \cdot y) \bmod 2^w)$$

- However, note that the **bit-level representation** of the product operation is **identical** for both unsigned and two's-complement multiplication

## H5 ¶2.3.6 Multiplying by Constants

- Historically, the **integer multiply instructions** on many machines was fairly slow
  - Other integer operations (+, -, bitwise, shifting) require  $\approx 1$  clock cycle, while multiplication takes  $\approx 3$  clock cycles even on the Intel Core i7 Haswell.
- To **optimize** multiplications, compilers replace multiplications by **constant factors** with **shift & addition** operations
  - e.g.  $x * 14$  can be rewritten as  $(x \ll 3) + (x \ll 2) + (x \ll 1)$  since  $14 = 2^3 + 2^2 + 2^1$



or,  $(x \ll 4) - (x \ll 1) \quad (\because 14 = 2^4 - 2^1)$

### Multiplication by a power of 2

Let  $x$  be the unsigned integer represented by bit pattern  $[x_{w-1}, x_{w-2}, \dots, x_0]$ .

Then for any  $k \geq 0$ , the  $w + k$ -bit unsigned representation of  $x2^k$  is given by  $[x_{w-1}, x_{w-2}, \dots, x_0, 0, \dots, 0]$ , where  $k$  zeros have been added to the right

When shifting left by  $k$  for a fixed word size, the high-order  $k$  bits are discarded, yielding

$$[x_{w-k-1}, x_{w-k-2}, \dots, x_0, 0, \dots, 0]$$

### Unsigned multiplication by a power of 2

For C variables  $x$  and  $k$  with unsigned values  $x$  and  $k$ , such that  $0 \leq k < w$ , the C expression  $x \ll k$  yields the value  $x *_{w}^u 2^k$

### Two's-complement multiplication by a power of 2

For C variables  $x$  and  $k$  with two's-complement value  $x$  and unsigned value  $k$ , such that  $0 \leq k < w$ , the C expression  $x \ll k$  yields the value  $x *_{w}^t 2^k$

## H5 ¶2.3.7 Dividing by Powers of 2

- **Integer division** is even slower than integer multiplication -->  $\approx 30$  or more clock cycles
  - Use **Right shifts!**
    - logical right shifts -- unsigned
    - arithmetic right shifts -- two's-complement
- Integer division always **rounds toward zero**

### Unsigned division by a power of 2

For C variables  $x$  and  $k$  with unsigned values  $x$  and  $k$ , such that  $0 \leq k < w$ , the C expression  $x \gg k$  yields the value  $\lfloor x/2^k \rfloor$ .

-- Note that for unsigned divisions, use **Logical Right Shifts!**

### Two's-complement division by a power of 2, rounding down

For C variables  $x$  and  $k$  have two's-complement value  $x$  and unsigned value  $k$ , respectively, such that  $0 \leq k < w$ . The C expression  $x \gg k$ , when the shift is performed **arithmetically**, yields the value  $\lfloor x/2^k \rfloor$

### Two's complement division by a power of 2, rounding up

Let C variables  $x$  and  $k$  have two's-complement value  $x$  and unsigned value  $k$ , respectively, such that  $0 \leq k < w$ . The C expression  $(x + (1 \ll k) - 1) \gg k$ , when the shift is performed **arithmetically**, yields the value  $\lceil x/2^k \rceil$

$$(\because \lceil x/y \rceil = \lfloor (x + y - 1)/y \rfloor)$$