

V. Optimizing Program Performance

- How to write efficient programs?
 1. appropriate set of algorithms and data structures
 2. write source code that the compiler can effectively optimize to turn into efficient executable code (CH.12)
 - need to understand the capabilities and limitations of optimizing compilers (pointer arithmetic, casting...)
 3. divide a task into portions that can be computed in parallel (multiple cores and multiple processors)
- To optimize written programs...
 1. eliminate unnecessary work (function calls, conditional tests, memory references)
 2. **instruction-level parallelism** (executing multiple instructions simultaneously)
 3. Use code **profilers** -- tools that measure the performance of different parts of a program

§5.1 Capabilities and Limitations of Optimizing Compilers

- Most compilers provide users with some control over **which optimizations** they apply
 - **Optimization level**
 - -Og : basic set of optimizations
 - -O1 or higher (-O2, -O3) causes more extensive optimizations
 - may expand the program size
 - may make the program harder to debug with standard debugging tools
- In CSAPP, will mostly consider code compiled with -O1
 - c.f. -O2 is standard for most software projects
- Compilers must be careful to apply only **safe** optimizations to a program
 - resulting program will have **the exact same behavior** as would an **unoptimized version** for all possible cases the program may encounter
- **Programmer** must make more of an effort to write program in a way that the **compiler can then transform into efficient machine-level code**

H6 Optimization Blockers

1. Memory aliasing

```
1 void twiddle1(long *xp, long *yp) {
2     *xp += *yp;
3     *xp += *yp;
4 }
5
6 void twiddle2(long *xp, long *yp) {
7     *xp += 2* *yp;
8 }
```

- both procedures seem to have identical behavior
- If yes, **twiddle2** is more efficient
 - twiddle2 - **3 memory references** (read *xp, read *yp, write *xp)
 - twiddle1 - **6 memory references** (2x read *xp, 2x read *yp, 2x write *xp)
- **However**, when **xp** and **yp** are equal
 - twiddle2 **triples value at xp**
 - twiddle1 **quadruples value at xp**
- Compiler **cannot generate** code in the style of **twiddle2** as an optimized version of twiddle1
- Memory aliasing - two pointers may designate the same memory location
 - Compiler **must assume** that **different pointers may be aliased**

```
1 x = 1000;   y = 3000;
2 *q = y;     // 3000
3 *p = x1;    // 1000
4 t1 = *q;    // 1000 or 3000
5             // 1000 if aliased
```

- Severely limit the opportunities for a compiler to generate optimized code

2. Function calls

```

1  long f();
2
3  long func1() {
4      return f() + f() + f() + f();
5  }
6
7  long func2() {
8      return 4*f();
9  }
10 // seems func1 and func2 compute the same result
11
12 /* What if f() does... */
13 long counter = 0;
14
15 long f() {
16     return counter++;
17 }

```

- Function **f** has a **side effect** -- modifies some part of the **global program state**
- Most compilers do not try to determine whether a function is **free of side effects**
--> leave function calls **intact**

c.f. **Optimizing function calls by inline substitution**

- function call is replaced by the code for the body of the function

```

1  long counter = 0;
2  /* Result of inlining f in func1 */
3  long func1lin() {
4      long t = counter++; // +0
5      t += counter++;     // +1
6      t += counter++;     // +2
7      t += counter++;     // +3
8      return t;
9  }
10
11 // after inlining, function can be optimized further
12 long func1opt() {
13     long t = 4 * counter + 6;
14     counter += 4;
15     return t;
16 } // reproduces the exact same behavior of func1

```

- recent versions of **GCC** attempts this form of optimization (when `-finline` or `-O1` or higher optimization level selected)
 - attempts inlining for functions defined within a single file
- When the code will be evaluated using a **symbolic debugger (GDB)**, better **not** to use **inline substitution**
 - tracing, setting breakpoints do not work properly

H3 §5.2 Expressing Program Performance

- **cycles per element (CPE)** will be used to express program performance in a way that can guide us in improving the code
 - helps us understand the **loop performance** of an iterative program at a detailed level
 - work well with programs that perform a **repetitive computation** (processing pixels or computing matrix products)
- but why **CPE**?
 - sequencing of activities is controlled by a clock providing a **regular signal of some frequency** (in GHz)
 - when the system uses "4GHz" processor, processor clock runs at 4.0×10^9 cycles per second
 - it is instructive for us programmers to use clock cycles rather than nanoseconds or picoseconds
 - focus more on how many instructions are being executed rather than how fast the clock runs
- Example code
 - functions **psum1** and **psum2** both compute the **prefix sum** of a vector of length **n**

For a vector $\vec{a} = \langle a_0, a_1, \dots, a_{n-1} \rangle$, the prefix sum $\vec{p} = \langle p_0, p_1, \dots, p_{n-1} \rangle$ is defined as

$$p_0 = a_0$$

$$p_i = p_{i-1} + a_i, \quad 1 \leq i < n$$

```
1  /* Compute prefix sum of vector a */
```

```

2 void psum1(float a[], float p[], long n) {
3     long i;
4     p[0] = a[0];
5     for (i = 1; i < n; i++)
6         p[i] = p[i-1] + a[i];
7 }
8
9 void psum2(float a[], float p[], long n) {
10    long i;
11    p[0] = a[0];
12    for (i = 1; i < n-1; i+=2) {
13        float mid_val = p[i-1] + a[i];
14        p[i] = mid_val;
15        p[i+1] = mid_val + a[i+1];
16    }
17    // For even n, finish remaining elem
18    if (i < n)
19        p[i] = p[i-1] + a[i];
20 }

```

- **psum2** uses an optimization technique called **loop unrolling** --> computing two elements per iteration
- time required by such a procedure can be characterized as a **Constant + a factor proportional to the number of elements**
 - e.g. **psum1** = $368 + 9.0n$, **psum2** = $368 + 6.0n$
 - > 368 cycles (timing code + initiating procedure, setting up the loop, completing the procedures) + linear factor of **6.0 or 9.0** cycles per element
 - > for large values of n, the run times will be dominated by the linear factors
 - **psum1** : CPE of 9.0, **psum2** : CPE of 6.0

H3 §5.3 Program Example

- To demonstrate how an abstract program can be systematically transformed into more efficient code, will use the **vector data structure**
- Header for defining vectors

```

1 // vec.h
2 /* Create abstract data type for vector */
3 typedef struct {
4     long len;
5     data_t *data;
6 } vec_rec, *vec_ptr;

```

- Generating vectors, accessing vector elements, and determining vector length

```

1 // vec.c
2 /* Create vector of specified length */
3 vec_ptr new_vec(long len) {
4     // Allocate header structure
5     vec_ptr result = (vec_ptr) malloc(sizeof(vec_rec));
6     data_t *data = NULL;
7     if(!result)
8         return NULL; // couldn't allocate storage
9     result->len = len;
10
11     // Allocate array
12     if(len > 0) {
13         data = (data_t *)calloc(len, sizeof(data_t));
14         if(!data) {
15             free((void *) result);
16             return NULL; // couldn't allocate storage
17         }
18     }
19     // Data will either be NULL or allocated arr
20     result->data = data;
21     return result;
22 }
23
24 /*
25  * Retrieve vector element and store at dest
26  * Return 0 (out of bounds) or 1 (successful)
27  */
28 int get_vec_element(vec_ptr v, long index, data_t *dest) {
29     if(index < 0 || index >= v->len)
30         return 0;
31     *dest = v->data[index];
32     return 1;
33 }

```

```

34
35  /* Return length of vector */
36  long vec_length(vec_ptr v) {
37      return v->len;
38  }

```

- Note that `get_vec_element` performs **bounds checking** for every vector reference
 - reduces the chance of program error but **can slow down** program execution
- Initial implementation of combining operation

```

1  /* Implementation with maximum use of data abstraction */
2
3  // Use one of the two definitions below
4  // For addition
5  #define IDENT 0
6  #define OP +
7
8  // For multiplication
9  #define IDENT 1
10 #define OP *
11
12 void combine1(vec_ptr v, data_t *dest) {
13     long i;
14
15     *dest = IDENT;
16     for(i = 0; i < vec_length(v); i++) {
17         data_t val;
18         get_vec_element(v, i, &val);
19         *dest = *dest OP val;
20     }
21 }

```

- Combines all of the elements in a vector into a single value according to some operation (addition or multiplication)
- **CPE** performance of the functions to be measured with an **Intel Core i7 Haswell processor** (CSAPP reference machine)
- CPE measurements for **combine1**

1			Integer		Floating pt	
2			-----			
3	Function	Method	+	*	+	*
4	-----					
5	combine1	unoptimized	22.68	20.02	19.98	20.18
6	combine1	Abstract -O1	10.12	10.12	10.17	11.14

--> with command-line option -O1, program performance significantly improved (more than 2x)

H3 §5.4 Eliminating Loop Inefficiencies

- Test condition **must be** evaluated on every iteration of the loop
 - in **combine1**, length of the vector does not change as the loop proceeds
 - compute the vector length only once and use this in test condition!

```

1  /* Move call to vec_length out of loop */
2  void combine2(vec_ptr v, data_t *dest) {
3      long i;
4      long length = vec_length(v);
5
6      *dest = IDENT;
7      for(i = 0; i < length; i++) {
8          data_t val;
9          get_vec_element(v, i, &val);
10         *dest = *dest OP val;
11     }
12 }
```

--> calls **vec_length** at the beginning and assigns the result to a local variable **length**

1				Integer		Floating pt	
2				-----			
3	Function	Method		+	*	+	*
4	-----						
5	combine1	Abstract -O1		10.12	10.12	10.17	11.14
6	combine2	move vec len		7.02	9.03	9.02	11.03

H6 Code Motion

- identifying a computation that is performed multiple times, **but such that the result of the computation will not change**
 - move the computation to an earlier section that does not get evaluated as often
 - Optimizing compilers attempt to perform code motion
 - but they assume that function has side effects
- > programmer must aid compiler by explicitly performing code motion

[Extreme example of Code motion]

```
1  /* Convert string to lowercase: slow */
2  void lower1(char *s) {
3      long i;
4      for (i = 0 ; i < strlen(s); i++)
5          if (s[i] >= 'A' && s[i] <= 'Z')
6              s[i] -= ('A' - 'a');
7  }
8
9  /* Convert string to lowercase: faster */
10 void lower2(char *s) {
11     long i;
12     long len = strlen(s);
13
14     for (i = 0; i < len; i++)
15         if (s[i] >= 'A' && s[i] <= 'Z')
16             s[i] -= ('A' - 'a');
17 }
18
19 /* Sample implementation of library function strlen */
20 /* Compute length of string */
21 size_t strlen(const char *s) {
22     long length = 0;
23     while (*s != '\0') {
24         s++;
25         length++;
26     }
27     return length;
28 }
```

- for a string of length n , **strlen** takes time proportional to n
 - since **strlen** is called in each of the n iterations of **lower1**, the overall run time of

lower1 is quadratic in the string length ==> proportional to n^2

- Code motion can become a major performance bottleneck

H3 §5.5 Reducing Procedure Calls

- **procedure calls** can incur overhead and block most forms of program optimizations
- in **combine2**, **get_vec_element** is called on every loop iteration to retrieve the next vector element
 - **get_vec_element** checks the vector index **i** against the loop bounds with every vector reference --> inefficiency
- Hence, suppose instead we add a function **get_vec_start** to our abstract data type

```
1  data_t *get_vec_start(vec_ptr v) {
2      return v->data;
3  }
4
5  /* Direct access to vector data */
6  void combine3(vec_ptr v, data_t *dest) {
7      long i;
8      long length = vec_length(v);
9      data_t *data = get_vec_start(v);
10
11     *dest = IDENT;
12     for(i = 0; i < length; i++) {
13         *dest = *dest OP data[i];
14     }
15 }
```

--> No function calls in the inner loop

--> Rather than making a function call to retrieve each vector element, it **accesses the array directly**

- However, in this case, there is **no apparent performance improvement**
 - but it is one of a series of steps that will **ultimately lead** to greatly improved performance

H3 §5.6 Eliminating Unneeded Memory References

- **combine3** accumulates the value being computed by the combining operation at the location designated by the pointer **dest**

[X86-64 code]

```
1  # Inner loop of combine3.  data_t = double, OP = *
2  # dest in %rbx, data+i in %rdx, data+length in %rax
3  .L17:                                #loop:
4      vmovsd  (%rbx), %xmm0            # Read pd from dest
5      vmulsd  (%rdx), %xmm0, %xmm0    # Multip pd by data[i]
6      vmovsd  %xmm0, (%rbx)           # Store pd at dest
7      addq    $8, %rdx                 # Increment data+i
8      cmpq    %rax, %rdx               # Compare to data+ln
9      jne     .L17                     # If !=, goto loop
```

- accumulated value is read from and written to memory on each iteration --> wasteful!
(since the value read from dest at the beginning of each iteration should simply be the value written at the end of the previous iteration)
- We can optimize the code by using a **temporary variable acc** that is used in the loop to accumulate the computed value
- [Combine 4 - use of temp var]

```
1  /* Accumulate result in local variable */
2  void combine4(vec_ptr v, data_t *dest) {
3      long i;
4      long length = vec_length(v);
5      data_t *data = get_vec_start(v);
6      data_t acc = IDENT;
7      /* Instead of using *dest, use temp var acc in loop */
8
9      for (i = 0; i < length; i++) {
10         acc = acc OP data[i];
11     }
12     *dest = acc;
13 }
```

- [Generated assembly code (**Inner loop**)]

```

1  # Inner loop of combine4, data_t = double, OP = *
2  # acc in %xmm0, data+i in %rdx, data+length in %rax
3  .L25:                                # Loop:
4      vmulsd    (%rdx), %xmm0, %xmm0    # Multip acc by data[i[]]
5      addq      $8, %rdx                 # Increment data+i
6      cmpq      %rax, %rdx               # Compare to data+length
7      jne       .L25                     # If !=, goto loop

```

			Integer		Floating pt	
			-----		-----	
	Function	Method	+	*	+	*

5	combine3	Direct access	7.17	9.02	9.02	11.03
6	combine4	acc in temp	1.27	3.01	3.01	5.01

H3 §5.7 Understanding Modern Processors

- to push the performance further, we must consider optimizations that exploit the microarchitecture of the processor
 - underlying system design by which a processor executes instructions
- general principles of operation and optimization apply to a wide variety of machines (yet, detailed performance results may not hold for other machines)
- at the **code level**, it appears as if instructions are **executed one at a time** (fetching values from registers or memory, performing an operation, and storing results back to a register or memory location)
- **BUT, in the ACTUAL PROCESSOR, a number of instructions are evaluated simultaneously --> Instruction-level parallelism**
 - can be 100 or more instructions run in parallel
- **Two different lower bounds** to maximum performance
 1. **Latency bound** : when a series of operations **must be** performed in **strict sequence**
 - when the result of one operation is required before the next one can begin

2. **Throughput bound:** raw computing capacity of the processor's functional units
(ULTIMATE limit on program performance)

H5 ¶5.7.1 Overall Operation

- recent **Intel processors** = **Superscalar**
 - can perform multiple operations on every clock cycle & out of order (don't need to correspond to their ordering in the machine-level program)
- **Overall Design** = **ICU** (Instruction control unit) + **EU** (execution unit)
 - **ICU** - reads a sequence of instructions from memory and generates a set of primitive operations to perform on program data
 - reads the instructions from an **instruction cache** (special high speed memory containing the most recently accessed instructions)
 - fetches **well ahead** of the currently executing instructions
 - **EU** - executes these operations

--> **out-of-order** processors require far greater & more complex hardware but they are better at achieving higher degrees of **instruction-level parallelism**

- When a program **hits a branch** -- two possible directions the program might go
 - Modern processors employ **Branch prediction**
 - guess whether or not a branch will be taken and predict the target address for the branch
 - begins executing these operations before it has been determined
 - If the **branch was predicted incorrectly**, **resets the state to that at the branch point** and begins fetching and executing instructions in the **other direction**
 - **Fetch control** (in ICU) incorporates branch prediction to perform the task of determining which instructions to fetch
- **Instruction decoding logic** takes the actual program instructions and converts them into a **set of primitive operations** (micro-operations)
 - each of these operations performs some simple computational task (adding two numbers, reading data from memory, writing data to memory)
 - for machines with **complex instructions** (x86 processors), instruction can be **decoded into multiple operations**
 - how instructions are decoded varies between machines
 - can **optimize** our programs without knowing the low-level details of a particular machine implementation

```
1  addq  %rax, %rdx
```

--> converted into a single operation

```
1  addq  %rax, 8(%rdx)
```

--> yields multiple operations, **separating the memory references** from the arithmetic operations

1. **load** a value from memory into the processor
2. **add** the loaded value to the value in register %eax
3. **store** the result back to memory

--> allow a division of labor among a set of dedicated hardware units (**multiple instructions in parallel**)

- **EU** (Execution Unit) then receives operations from the instruction fetch unit
 - receives a number of operations on each clock cycle
 - operations are dispatched to a set of **functional units** (specialized to handle different types of operations)

1. **Load unit** - **reads** data from the memory into the processor
 - has an **adder** to perform **address computations**
2. **Store unit** - **writes** data from the processor to the memory
 - has an **adder** to perform **address computations**

--> Load unit & Store unit access memory via **data cache** (high-speed memory containing the most recently accessed data values)

3. **Arithmetic operations unit**

- typically **specialized to perform** different combinations of **integer & floating-point operations**
- as # of transistors ↑↑

=> total # of functional units & combinations of operations each unit can perform & performance of each of these units ↑↑

--> Designed to be able to perform a variety of different operations (If one functional unit were specialized to perform specific operation, would not get the full benefit of having multiple functional units)

4. **Branch unit**

- **final results are not stored in the program registers or data memory until the processor can be certain that these instructions should actually have been executed**
- **determines** whether or not **branch operations were predicted correctly**
 - If the prediction was **incorrect**, **EU discards** the results that have been computed **beyond the branch point** & **signals the branch unit** that the prediction was incorrect and **indicate the correct branch destination**
- **Intel Core i7 Haswell** has eight functional units (numbered 0-7) -- partial list of each one's capabilities --
 0. Integer arithmetic, floating-point multiplication, integer and floating-point division, branches
 1. Integer arithmetic, floating-point addition, integer multiplication, floating-point multiplication
 2. Load, address computation
 3. Load, address computation
 4. Store
 5. Integer arithmetic
 6. Integer arithmetic, branches
 7. Store address computation

--> Combination of functional units has the potential to perform **multiple operations of the same type simultaneously** (∴ 4 units for integer operations, 2 units for load operations, 2 for floating-point multiplication)

- **Retirement unit** keeps track of the ongoing processing and makes sure that it obeys the sequential semantics of the machine-level program
 - controls the updating of **register files**
 - as an instruction is decoded, information about instruction is placed into a **FIFO queue**, and remains until...
 1. operations for the instructions have completed and any branch points leading to this instruction are confirmed as having been **correctly predicted** --> instruction can be **retired**, with any updates to the program registers
 2. If **mispredicted**, instruction will be **flushed**, discarding any results that may have been computed --> **will not alter the program state**
- to **expedite the communication of results** from one instruction to another, information is **exchanged among the execution units** (operation results)

- execution units can send results directly to each other
- Register renaming
 - when an instruction that updates **register r** is decoded, a **tag t** is generated, giving a unique identifier to the result of the operation
 - **(r, t)** is added to a **table** maintaining the **association between program register r and tag t** for an operation that will update this register
 - when some execution unit **completes** the first operation, it generates a result **(v, t)**, indicating that the operation with **tag t produced value v**
 - any operation waiting for **t** as a source will then use **v as the source value (data forwarding)**
 - values can be forwarded directly from one operation to another (without registers), enabling the second operation to begin as soon as the first has completed
 - with **register renaming**, an entire sequence of operations can be performed **speculatively**, even though the registers are updated only after the processor is certain of the branch outcomes (∴ information is passed directly within execution unit)

H5 ¶5.7.2 Functional Unit Performance

[Performance of arithmetic operations for i7 Haswell]

1		Integer			FP		
2		-----					
3	Op	Latency	Issue	Capac	Latency	Issue	Capac
4		-----					
5	+	1	1	4	3	1	1
6	*	3	1	1	5	1	2
7	/	3-30	3-30	1	3-15	3-15	1
8		-----					

- **Latency:** total time required to perform the operation
 - latencies increase in going from integer to floating-point operations
- **Issue time:** minimum number of clock cycles between two independent operations of the same type
 - addition & multiplication operations all have issue time of 1

- on each clock cycle, the processor can start a new one of these operations
- achieved through the use of **pipelining**
 - **pipelined function unit** is implemented as a series of **stages** (performs part of the operation)

e.g. floating-point adder contains 3 stages (3-cycle latency)

 1. process the exponent values
 2. add the fractions
 3. round the result
 - arithmetic operations can proceed through the stages in close succession (rather than waiting for one operation to complete before the next begins)
 - only if there are successive, logically independent operations to be performed
- functional units with issue times of 1 cycle are said to be **fully pipelined** --> can start a new operation every clock cycle
- **divider is not pipelined** --> issue time = latency
 - divider **must perform** a complete division before it can begin a new one
 - comparatively costly operation
- **Capacity:** number of functional units capable of performing the operation
 - operations with capacity greater than 1 = due to the capabilities of the **multiple functional units**
- **Throughput** of the unit?
 - more common way of expressing issue time = specifying the maximum **throughput** of the unit
 - **reciprocal of the issue time**
 - **fully pipelined functional** unit has a maximum throughput of **1 operation per clock cycle**, while **units with higher issue times** have **lower maximum throughput**
 - having **multiple functional units can increase throughput**

operation with capacity C and issue time I can potentially be achieved a throughput of C/I operations per clock cycle
- CPU designers must carefully balance the number of functional units and their individual performance to achieve optimal overall performance (∴ limited amount of space on the microprocessor chip)

- in the design of the Core i7 Haswell processor...
 - integer multiplication and floating-point multiplication and addition were considered important operations
 - > most resources dedicated
 - division is relatively infrequent and difficult to implement with either short latency or full pipelining
- **latency bound** -- minimum value for the CPE for any function that must perform the combining operation in a strict sequence
- **throughput bound** -- minimum bound for the CPE based on the maximum rate at which the functional units can produce results

H5 ¶5.7.3 An Abstract Model of Processor Operation

- to analyze the performance of a machine-level program executing on a modern processor, we use **data-flow** representation of programs, showing **data dependencies** between the different operations constrain the order in which they are executed
 - constraints lead to critical path in the graph, putting a lower bound on the number of clock cycles required to execute a set of machine instructions

c.f. computing the product or sum of n elements requires

$\approx L \cdot n + k$ clock cycles where L is the latency of the combining operation and K represents the overhead of calling the function and initiating and terminating the loop

==> CPE is equal to the latency bound L

- For a code segment forming a **loop**, we can classify the registers that are accessed into **four categories**
 1. Read-only
 - used as source values (as data or to compute memory addresses)
 - **not modified within the loop**
 2. Write-only
 - used as the **destinations** of data-movement operations
 3. Local

- updated and used within the loop, but there is **no dependency from one iteration to another**
- e.g. **Condition code registers**

4. Loop

- used both as **source values** and as **destinations** for the loop
 - value generated in one iteration being used in another
-
- **For all of the cases where the operation has a latency L greater than 1, we see that the measured CPE is simply L , indicating that this chain forms the performance-limiting critical path**
-
- critical paths in a **data-flow representation** provide **lower bound** on how many cycles a program will require
 - **other factors can also limit performance** (total number of functional units available & number of data values that can be passed among the functional units on any given step)
 - e.g. **Integer addition** -- data operation is sufficiently fast that the rest of the operations cannot supply data fast enough

H3 §5.8 Loop Unrolling

- program transformation that reduces the number of iterations for a loop by **increasing the number of elements computed on each iteration**
- loop unrolling can improve performance in two ways...
 1. **reduces the number of operations that do not contribute directly** to the program result
 - loop indexing / conditional branching
 2. **exposes** ways in which we can further transform the code to **reduce the number of operations in the critical paths** of the overall computation

[Example C code, **2 x 1 loop unrolling**]

```
1  /* 2 x 1 loop unrolling */
2  void combine5(vec_ptr v, data_t *dest) {
3      long i;
4      long length = vec_length(v);
5      long limit = length-1;
```

```

6     data_t *data = get_vec_start(v);
7     data_t acc = IDENT;
8
9     /* Combine 2 elements at a time */
10    for (i = 0; i < limit; i+=2) {
11        acc = (acc OP data[i]) OP data[i+1];
12    }
13
14    /* Finish any remaining elements */
15    for (; i < length; i++) {
16        acc = acc OP data[i];
17    }
18    *dest = acc;
19 }

```

--> first loop steps through the array **two elements at a time** , combining operation is applied to array elements i and $i + 1$ in a single iteration

--> since the vector length will not always be a multiple of 2, second loop operates on the remaining element(s)

--> can generalize the idea to **unroll a loop by any factor k** --> yielding $k \times 1$ **loop unrolling** (unroll by a factor of k but accumulate values in a single variable acc)

1) set the upper limit to be $n - k + 1$, within the loop apply the combining operation to elements i through $i + k - 1$

2) loop index i is incremented by k in each iteration

3) maximum array index $i + k - 1$ will be less than n

4) include the second loop to step through the final few elements of the vector one at a time

[Performance of unrolled code for unrolling factors $k = 2$ and $k = 3$]

1			Integer		Floating pt	
2			-----			
3	Function	Method	+	*	+	*
4	-----					
5	combine4	No unrolling	1.27	3.01	3.01	5.01
6	combine5	2x1 unrolling	1.01	3.01	3.01	5.01
7		3x1 unrolling	1.01	3.01	3.01	5.01
8						
9	Latency bound		1.00	3.00	3.00	5.00
10	Throughput bound		0.50	1.00	1.00	0.50

- CPE for **integer addition improves**, achieving the **latency bound of 1.00**
 - benefits of **reducing loop overhead operations**
 - reducing the number of overhead operations relative to the number of additions required to compute the vector sum, **can reach the point where 1-cycle latency of integer addition becomes the performance-limiting factor**
 - **none go below their latency bounds** (even for bigger unrolling factors)
 - ($\therefore 1/k$ iterations of nk operations per iteration \rightarrow critical path of n operations is the **limiting factor** for the performance)

[Assembly code for the inner loop of combine5 ($k = 2$, data_t = double)]

```

1  # Inner loop of combine5. data_t = double, OP = *
2  # i in %rdx, data %rax, limit in %rbx, acc in %xmm0
3  .L35:                                # loop:
4      vmulsd  (%rax,%rdx,8), %xmm0, %xmm0  # Multiply acc by data[i]
5      vmulsd  8(%rax,%rdx,8), %xmm0, %xmm0  # Multiply acc by data[i+1]
6      addq    $2, %rdx                    # Increment i by 2
7      cmpq    %rdx, %rbp                  # Compare to limit:i
8      jg      .L35                        # If >, goto loop

```

H3 §5.9 Enhancing Parallelism

- **functional units** performing **addition** and **multiplication** are fully pipelined \rightarrow can start **new operations every clock cycle**
- some operations can be performed by **multiple functional units**

--> has the potential to perform additions and multiplications at a higher rate

--> but in **combine5**, we used a single variable **acc** to accumulate the values. If we use more variables...?

H5 ¶5.9.1 Multiple Accumulators

- improve performance by **splitting the set of combining operations into two or more parts and combining the results at the end**

Let P_n denote the product of elements a_0, a_1, \dots, a_{n-1} :

$$P_n = \prod_{i=0}^{n-1} a_i$$

Assuming n is even, we can write this as $P_n = PE_n \times PO_n$, where PE_n is the product of the elements with even indices, and PO_n is the product of the elements with odd indices:

$$PE_n = \prod_{i=0}^{n/2-1} a_{2i}$$

$$PO_n = \prod_{i=0}^{n/2-1} a_{2i+1}$$

[Example C code, **2 x 2 loop unrolling**]

```
1  /* 2 x 2 loop unrolling */
2  void combine6(vec_ptr v, data_t *dest) {
3      long i;
4      long length = vec_length(v);
5      long limit = length-1;
6      data_t *data = get_vec_start(v);
7      data_t acc0 = IDENT;
8      data_t acc1 = IDENT;
9
10     /* Combine 2 elements at a time */
11     for (i = 0; i < limit; i += 2) {
12         acc0 = acc0 OP data[i];
13         acc1 = acc1 OP data[i+1];
14     }
15
16     /* Finish any remaining elements */
17     for (; i < length; i++) {
18         acc0 = acc0 OP data[i];
19     }
20     *dest = acc0 OP acc1;
21 }
```

- uses **two-way loop unrolling** and **two-way parallelism**

[Performance comparison]

1			Integer		Floating pt	
2			-----			
3	Function	Method	+	*	+	*
4	-----					
5	combine4	No unrolling	1.27	3.01	3.01	5.01
6	combine5	2x1 unrolling	1.01	3.01	3.01	5.01
7	combine6	2x2 unrolling	0.81	1.51	1.51	2.51
8						
9	Latency bound		1.00	3.00	3.00	5.00
10	Throughput bound		0.50	1.00	1.00	0.50

- improved the performance for all cases by a **factor of around 2**
- **broken through the barrier** imposed by the **latency bound**
 - **two critical paths** - one for even-numbered elements, and one for the odd-numbered elements
 - each contains only $n/2$ operations --> CPE /2
 - only **exception** = **integer addition**
 - too much loop overhead to achieve the theoretical limit of 0.50
- in general, a program can **achieve the throughput bound** for an operation only when it can **keep the pipelines filled for all of the functional units capable of performing that operation**
 - for an operation with latency L and capacity C , this requires an unrolling factor $k \geq C \cdot L$
- **MUST consider** whether it **preserves the functionality** of the original function
 - **two's-complement arithmetic** is **commutative** and **associative** even when **overflow occurs**
 - optimizing compiler could convert the code by loop unrolling & parallelism
 - **floating-point multiplication and addition** are **not associative**
 - loop unrolling & parallelism could produce different results due to **rounding** or **overflow**
 - however, **achieving a performance gain of 2x outweighs the risk of generating different results** for strange data patterns

- in real-life applications it is unlikely that multiplying the elements in strict order gives fundamentally better accuracy than does multiplying two groups independently and then multiplying those products together
- most compilers **do not attempt** such transformations since they have no way to judge the risks of introducing transformations can change the program behavior (no matter how small it is)

H5 ¶5.9.2 Reassociation Transformation

[Example C code, 2 x 1a loop unrolling]

```

1  /* 2 x 1a loop unrolling */
2  void combine7(vec_ptr v, data_t *dest) {
3      long i;
4      long length = vec_length(v);
5      long limit = length-1;
6      data_t *data = get_vec_start(v);
7      data_t acc = IDENT;
8
9      /* Combine 2 elements at a time */
10     for (i = 0; i < limit; i+=2) {
11         acc = acc OP (data[i] OP data[i+1]);
12     }
13
14     /* Finish any remaining elements */
15     for (; i < length; i++) {
16         acc = acc OP data[i];
17     }
18     *dest = acc;
19 }
```

[Difference between combine5 - combine7]

```

1  /* combine5 */
2  acc = (acc OP data[i]) OP data[i+1];
3
4  /* combine7 */
5  acc = acc OP (data[i] OP data[i+1]);
```


[Performance Comparison]

1			Integer		Floating pt	
2			-----			
3	Function	Method	+	*	+	*
4	-----					
5	combine4	No unrolling	1.27	3.01	3.01	5.01
6	combine5	2x1 unrolling	1.01	3.01	3.01	5.01
7	combine6	2x2 unrolling	0.81	1.51	1.51	2.51
8	combine7	2x1a unrollin	1.01	1.51	1.51	2.51
9						
10	Latency bound		1.00	3.00	3.00	5.00
11	Throughput bound		0.50	1.00	1.00	0.50

- except for integer addition, other **three cases double the performance** relative to $k \times 1$ unrolling, **breaking** through the barrier imposed by the **latency bound**
 - single critical path with only $n/2$ operations**
- Reassociation transformation** can **reduce the number of operations along the critical path in a computation**, resulting in better performance by **better utilizing the multiple functional units and their pipelining capabilities**
 - However, in general, **loop unrolling + accumulating multiple values in parallel** is more reliable way to achieve improved program performance

H3 §5.10 Summary of Results for Optimizing Combining Code

1			Integer		Floating pt	
2			-----			
3	Function	Method	+	*	+	*
4	-----					
5	combine1	Abstract -O1	10.12	10.12	10.17	11.14
6	combine6	2x2 unrolling	0.81	1.51	1.51	2.51
7		10x10 unrolling	0.55	1.00	1.01	0.52
8						
9	Latency bound		1.00	3.00	3.00	5.00
10	Throughput bound		0.50	1.00	1.00	0.50

- all been done using **ordinary C code and a standard compiler**
- by rewriting the code to take advantage of the newer **SIMD** instructions yields **additional performance gains of $\approx 4x$ or $8x$**
 - for floating point multiplication, CPE drops to 0.06 (180x performance)

H3 §5.11 Some Limiting Factors

- **Major limiting factors so far**
 1. **critical path** in a data flow
 - if there is **some chain of data dependencies** (sum of all of the latencies = T), then the program will require at least T cycles to execute
 2. **throughput bounds** of the functional units
 - assuming that a program requires a total of N computations of some operation, and the microprocessor has C functional units capable of performing that operation, and have an issue time of I , then the program will require at least $N \cdot I/C$ cycles to execute
- **Two Additional** limiting factors...

H5 ¶5.11.1 Register Spilling

- benefits of loop parallelism are limited by the **ability to express the computation in assembly code**
- If a program has a degree of parallelism P that **exceeds the number of available registers**, then the compiler will resort to **spilling**
 - **storing** some of the **temporary values in memory** by allocating space on the **run-time stack**

1			Integer		Floating pt	
2			-----		-----	
3	Function	Method	+	*	+	*
4	-----					
5	combine6					
6		10x10 unrolling	0.55	1.00	1.01	0.52
7		20x20 unrolling	0.83	1.03	1.02	0.68
8						
9	Throughput	bound	0.50	1.00	1.00	0.50

- modern **x86-64** processors have **16 integer registers** and can make use of the **16 YMM registers** to store floating-point data
 - once the number of loop variables **exceeds** the number of available registers --> **program must allocate some on the stack**

```

1  # Updating of accumulator acc0 in 10x10 unrolling
2  vmulsd  (%rdx), %xmm0, %xmm0    # acc0 *= data[i]
3
4  # Updating of accumulator acc0 in 20x20 unrolling
5  vmovsd  40(%rsp), %xmm0
6  vmulsd  (%rdx), %xmm0, %xmm0
7  vmovsd  %xmm0, 40(%rsp)

```

- the **accumulator** is kept as a **local variable on the stack**
 - program **MUST read** both its value and the value of data[i] from memory, **multiply** them and **store** the result back to memory

H5 ¶5.11.2 Branch Prediction and Misprediction Penalties

- conditional branch can incur a **significant misprediction penalty** when the branch prediction logic does not correctly anticipate whether or not a branch will be taken
- modern processors work well ahead of the currently executing instructions (**instruction pipelining**)
 - when a **branch is encountered**, the processor **must guess** which way the branch will go
- with **speculative execution**, the processor **begins executing** the instructions at the **predicted branch target**
 - **while not modifying any actual register or memory locations** until the actual outcome has been determined
 - if the prediction is **correct** --> processor **"commits"** the results of the speculatively executed instructions by **storing** them in registers or memory
 - if the prediction is **incorrect** --> processor **"must discard"** all of the speculatively executed results and **restart** the instruction fetch process at the correct location
- **Conditional move** instructions (rather than conditional transfers of control)
 - compute the values along both branches of a conditional expression or statement and then **use conditional moves** to **select the desired value**
 - no need to guess whether or not the condition will hold --> **no penalty for guessing incorrectly**

H6 Do Not Be Overly Concerned about Predictable Branches

- **branch prediction logic** found in modern processors is very good at **discerning regular patterns and long-term trends** for the different branch instructions
- performance of predicted branches is **limited by the latencies of operations**
 - processor is able to **predict** the outcomes of branches, so **none of this evaluation has much effect** on the fetching and processing of the **instructions that form the critical path** in the program execution

H6 Write Code Suitable for Implementation with Conditional Moves

- **branch prediction is only reliable for regular patterns**
 - many tests are completely **unpredictable!**
 - program performance can be **greatly enhanced** if the compiler is able to generate code **using conditional data transfers** rather than conditional control transfers

[Example 'Naive' C code]

```
1  /* Rearrange two vectors, for each i, b[i] >= a[i] */
2  void minmax1(long a[], long b[], long n) {
3      long i;
4      for (i = 0; i < n; i++) {
5          if (a[i] > b[i]) {
6              long t = a[i];
7              a[i] = b[i];
8              b[i] = t;
9          }
10     }
11 }
```

--> CPE of ≈ 13.5 for random data & ≈ 2.5 -3.5 for predictable data

[Example 'Better' C code - conditional data transfers]

```

1  /* Rearrange two vectors, for each i, b[i] >= a[i] */
2  void minmax2(long a[], long b[], long n) {
3      long i;
4      for (i = 0; i < n; i++) {
5          long min = a[i] < b[i] ? a[i] : b[i];
6          long max = a[i] < b[i] ? b[i] : a[i];
7          a[i] = min;
8          b[i] = max;
9      }
10 }

```

--> CPE of ≈ 4.0 regardless of whether the data are arbitrary or predictable

- not all conditional behavior can be implemented with conditional data transfers
 - \exists cases where programmers cannot avoid writing code that will lead to conditional branches for which the processor will do poorly with its branch predictions
 - try best to use **conditional data transfers**

H3 §5.12 Understanding Memory Performance

- all modern processors contain one or more cache memories to provide **fast access** to such small amounts of memory
- modern processors have dedicated **functional units** to perform load and store operations
 - these units have **internal buffers** to hold sets of outstanding requests for memory operations
 - i7 Haswell has... (can initiate 1 operation every clock cycle)
 - **2x load units** - each can hold up to **72** pending read requests
 - **1x store unit** - can hold up to **42** write requests

H5 ¶5.12.1 Load Performance

- performance of a program containing load operations depends on **1) pipelining capability** and **2) latency of the load unit**

- 1) one factor limiting the CPE is the fact that they all require **reading one value from memory** for each element computed
 - with two load units (each able to initiate at most 1 load operation every clock cycle), CPE cannot be less than 0.50
 - for applications loading k values for every element computed, CPE cannot be lower than $k/2$
- 2) performance effects due to the **latency of load operations**

[Example C code]

```

1  typedef struct ELE {
2      struct ELE *next;
3      long data;
4  } list_ele, *list_ptr;
5
6  long list_len(list_ptr ls) {
7      long len = 0;
8      while (ls) {
9          len++;
10         ls = ls->next;
11     }
12     return len;
13 }

```

--> sequence of load operations (outcome of one determines the address for the next)

--> **list_len** has a CPE of 4.00

[Generated **Assembly code**]

```

1  # Inner loop of list_len
2  # ls in %rdi, len in %rax
3  .L3:                                # loop:
4      addq    $1, %rax                # Increment len
5      movq    (%rdi), %rdi            # ls = ls->next
6      testq   %rdi, %rdi              # Test ls
7      jne     .L3                    # If nonnull, goto loop

```

--> **movq** instruction on line 5 forms the **critical bottleneck** in this loop

--> each successive value of register %rdi depends on the result of a load operation having the value in %rdi as its address

--> load operation for one iteration cannot begin until the one for the previous iteration has completed

H5 ¶5.12.2 Store Performance

- **store** operation -- writes a register value to memory
 - can operate in a **fully pipelined mode**, beginning a new store on every cycle (CPE of 1.0, with a single store functional unit)
 - **DOES NOT** affect any register values
 - series of store operations **CANNOT** create a **data dependency**
- **Store unit** includes a **store buffer** containing the **addresses and data of the store operations** that have been issued to the store unit, **but, have not yet been completed** (completion = updating the data cache)
 - series of store operations can be executed without having to wait for each one to update the cache
 - when a **load operation occurs**, it **must check the entries in the store buffer** for matching addresses
 - If it finds a match (any of the bytes being written have the same address as any of the bytes being read), it **retrieves the corresponding data entry** as the result of the load operation

[Example C code]

```
1  /* Set elements of array to 0 */
2  void clear_array(long *dest, long n) {
3      long i;
4      for (i = 0; i < n; i++)
5          dest[i] = 0;
6  }
7
8  /* Write to dest, read from src */
9  void write_read(long *src, long *dst, long n) {
10     long cnt = n;
11     long val = 0;
12
13     while (cnt) {
14         *dst = val;
15         val = (*src)+1;
```

```

16     cnt--;
17 }
18 }

```

--> if function **write_read** is called with arguments **src** and **dest** pointing to the **same memory location**, **write/read dependency** causes a slowdown in the processing

[Generated **Assembly code**]

```

1  # Inner loop of write_read
2  # src in %rdi, dst in %rsi, val in %rax
3  .L3:                                # loop:
4  movq    %rax, (%rsi)                # Write val to dst
5  movq    (%rdi), %rax                # t = *src
6  addq    $1, %rax                    # val = t+1
7  subq    $1, %rdx                    # cnt--
8  jne     .L3                        # If != 0, goto loop

```

- With matching source and destination addresses, **data dependency** between the **s_data** (movq instruction in L4 is translated into two operations: 1) s_addr, 2) s_data where s_addr computes the address for the store operation, creates an entry in the store buffer, and sets the address field for that entry, and s_data sets the data field for the entry) causes a **critical path** to form involving data being stored, loaded, and incremented
- with operations on **registers** -- **processor can determine** which instructions will affect which others as they are being decoded into operations
- with operations on **memories** -- **processor cannot predict** which will affect which others until the load and store addresses have been computed

H3 §5.13 Life in the Real World: Performance Improvement Techniques

Basic strategies for optimizing program performance:

1. **High-level design** -- choose appropriate algorithms and data structures for the problem at hand

2. **Basic coding principles** -- avoid optimization blockers so that a compiler can generate efficient code
 - Eliminate excessive function calls
 - move computations out of loops when possible
 - Eliminate unnecessary memory references
 - introduce temporary variables to hold intermediate results
 - store a result in an array or global variable only when the final value has been computed
3. **Low-level optimizations** -- structure code to take advantage of the hardware capabilities
 - **unroll loops** to reduce overhead and to enable further optimizations
 - find ways to increase **instruction-level parallelism** by techniques such as **multiple accumulators** and **reassociations**
 - rewrite conditional operations in a functional style to enable compilation via **conditional data transfers**

H3 §5.14 Identifying and Eliminating Performance Bottlenecks

- when working with large programs, it may be difficult to find where to focus more for optimization
- **code profilers** ; analysis tools that collect **performance data** about a program as it executes

H5 ¶5.14.1 Program Profiling

- to determine how much time the different parts of the program require
- Unix systems provide the profiling program **GPROF**
 1. determines how much **CPU time** was spent for each of the functions in the program
 2. computes a **count of how many times each function gets called** , categorized by which function performs the call
- Profiling with **GPROF** in three steps

Suppose we have a C program **prog.c** which runs with command-line argument **file.txt**

1. program must be compiled and linked for profiling
 - include the run-time flag **-pg** on the command line

```
1 linux> gcc -Og -pg prog.c -o prog
```

2. program is then executed as usual

```
1 linux> ./prog file.txt
```

- runs slightly slower than normal, generating a file **gmon.out**

3. GPROF is invoked to analyze the data in **gmon.out**

```
1 linux> gprof prog
```

- the **first part** of the profile report **lists the times spent executing the different functions** , sorted in descending order
- the **second part** of the profile report shows the **calling history of the functions**
- Properties of **GPROF**
 1. timing is **not very precise**
 2. calling information is **quite reliable**
 3. by default, **the timings for library functions are not shown** , these times are **incorporated into the times for the calling functions**

H3 §5.15 Summary

- No compiler can replace an inefficient algorithm or data structure by a good one
--> **better program design** should be a primary concern for programmers

1. **Eliminate optimization blockers**

- memory aliasing
- procedure calls

2. **Know processor's microarchitecture** (OOO processors)
 - capabilities, latencies, issue times of the functional units
3. **Exploit the instruction-level parallelism** provided by modern processors
 - loop unrolling
 - creating multiple accumulators
 - reassociation
4. **Analyze data dependencies**
 - critical paths between the different iterations of a loop
5. **Use Conditional data transfers**
 - make branches more predictable
 - keep values in local variables, allowing them to be stored in registers
6. **Use Code profilers**
 - to analyze the program