

## VIII. Exceptional Control Flow

---

- While power is on, the **Program counter** assumes a sequence of values  $a_0, a_1, \dots, a_{n-1}$  where  $a_k$  is the address of some corresponding instruction  $I_k$
- **Control transfer** -- each transition from  $a_k$  to  $a_{k+1}$
- **Control flow** -- sequence of such control transfers
  - "smooth" sequence when each  $I_k$  and  $I_{k+1}$  are adjacent in memory
  - abrupt changes caused by **jumps, calls, returns**
    - necessary mechanisms that allow programs to react to changes in **internal program state** represented by program variables
- **Systems must** also be able to **react to changes** in system state that are **not captured by internal program variables** and are **not necessarily related to the execution of the program**
  - Examples
    - hardware timer goes off at regular intervals
    - packets arrive at the network adapter and must be stored in memory
    - programs request data from a disk and sleep until they are notified that the data are ready
    - parent processes that create child processes must be notified when their children terminate
- Modern systems react to abovementioned situations by making abrupt changes (**exceptional control flow (ECF)**) in control flow
  - **ECF** occurs at all levels of a computer system
    - **Hardware** level -- events detected by the hardware trigger abrupt control transfers to **exception handlers**
    - **Operating systems** level -- kernel transfers control from one user process to another via **context switches**
    - **Application** level -- process can send a **signal** to another process that abruptly transfers control to a **signal handler**
    - **Individual program** -- sidestepping usual stack discipline and making **nonlocal jumps** to arbitrary locations in other functions
- Importance of **ECF** (Understanding ECF will help you...)

1. understand important systems concepts
  - OS uses **ECF** to implement I/O, processes, and virtual memory
2. understand how applications interact with the operating system
  - applications request services from the OS by using a form of ECF known as a **trap** or **system call**
    - writing data to a disk, reading data from a network, creating a new process, terminating the current process
3. write interesting new application programs
  - Unix shells, Web servers
4. understand concurrency
  - exception handler interrupts the execution of an application program
  - processes and threads whose execution overlap in time
  - signal handler that interrupts the execution of an application program
5. understand how software exceptions work
  - software exceptions allow the program to make **nonlocal** jumps in response to error conditions

### H3 §8.1 Exceptions

- **Exceptions** -- a form of exceptional control flow that are implemented partly by the hardware and partly by the operating system
  - abrupt change in the control flow **in response to** some **change in the processor's state**
- Suppose the processor is executing some current instruction  $I_{curr}$  when a significant change (**event**) in the processor's **state** occurs
  - state -- encoded in various bits and signals inside the processor
  - **event** might be directly related to the execution of the current instruction
- When the processor **detects** that the **event has occurred**, it makes an **indirect procedure call (exception)**, through a jump table called an **exception table** to an operating subroutine (**exception handler**)
  - when the exception handler finishes processing, **three** possibilities (depends on the type of event that caused the exception)
    1. handler returns control to the current instruction  $I_{curr}$ , the instruction that was executing when the event occurred
    2. handler returns control to  $I_{next}$ , the instruction that would've executed next

- had the exception not occurred
3. handler aborts the interrupted program

## H5 ¶8.1.1 Exception Handling

- Each type of possible exception in a system is assigned a **unique nonnegative integer exception number**
  - assigned by the designers of the **processor**
    - divide by zero, page faults, memory access violations, break points, arithmetic overflows
  - assigned by the designers of the **operating system kernel** (memory-resident part of the OS)
    - system calls, signals from external I/O devices
- At system boot time (reset or powered on), the **OS** allocates and initializes an **exception table**
  - entry  $k$  contains the **address of the handler for exception  $k$**
- At run time, the **processor** detects that an **event has occurred** and determines the corresponding **exception number  $k$** 
  - then triggers the exception by making an **indirect procedure call**, through entry  $k$  of the exception table, to the **corresponding handler**
  - exception number = index into the exception table
    - starting address contained in a special CPU register ( **exception table base register** )
- Some remarkable properties of **exceptions**
  - return address is **either the current instruction or the next instruction**
  - processor **pushes some additional processor state onto the stack** that will be necessary to **restart the interrupted program** when the handler returns
  - when control is being transferred **from a user program to the kernel**, all of the items are **pushed onto the kernel's stack**
  - **exception handlers** run in **kernel mode** (complete access to all system resources)
  - after the handler has processed the event, it **optionally** returns to the interrupted program by executing "**return from interrupt**" instruction which
    1. pops the appropriate stack back into the processor's control and data registers

2. restores the state to **user mode** (if the exception interrupted a user program)
3. returns control to the interrupted program

## H5 ¶8.1.2 Classes of Exceptions

- Exceptions can be divided into four classes
  1. **Interrupts**
  2. **Traps**
  3. **Faults**
  4. **Aborts**

1	Class	Cause	Async/Sync	Return
	behavior			
2	Interrupt	Signal from I/O device	Async	Always to
	next instruction			
3	Trap	Intentional exception	Sync	Always to
	next instruction			
4	Fault	Potentially recoverable error	Sync	Might reurun
	to current instr			
5	Abort	Nonrecoverable error	Sync	Never
	returns			

## H6 Interrupts

- Occur **asynchronously** as a result of **signals** from I/O devices that are external to the processor
  - **not caused by** the execution of any particular instruction
  - handled by **interrupt handlers**
- Brief process for an interrupt
  - I/O devices(network adapters, disk controllers, timer chips...) **trigger interrupts by signaling a pin** on the processor chip and **placing on to the system bus the exception number** that identifies the device that caused the interrupt
  - After the current instruction finishes executing, the **processor notices** that the interrupt pin has gone high, **reads the exception number** from the system bus, then **calls the appropriate interrupt handler**
  - When the handler returns, it **returns control to the next instruction**

--> **program continues executing as though the interrupt had never happened**

## H6 Traps and System Calls

- Traps -- **intentional** exceptions that occur as a result of executing an instruction
  - **trap handlers** return control to the next instruction
  - provide a procedure-like interface **between user programs and the kernel** --> system call
- User programs **request services from the kernel** when...
  1. reading a file ( **read** )
  2. creating a new process ( **fork** )
  3. loading a new program ( **execve** )
  4. terminating the current process ( **exit** )
- Processors provide a special **syscall**  $n$  instruction that user programs **can execute** when they want to request service  $n$ 
  - executing a **syscall** causes a **trap** to an **exception handler** that decodes the argument and **calls the appropriate kernel routine**
- From a programmer's perspective, a system call is identical to a regular function call, **However**, a **system call** runs in kernel mode (which allows it to **execute privileged instructions & access a stack defined in the kernel** )

## H6 Faults

- Result from error conditions that a **handler might be able to correct**
- When a fault occurs, the **processor transfers control to the fault handler**
  - If the **handler is able to correct the error condition** , it returns control to the **faulting instruction** (re-execution)
  - If the handler **cannot** correct the error condition, **returns to an abort routine** in the kernel, **terminating the application program** that caused the fault
- Example) **Page fault exception**
  - when an instruction references a virtual address whose corresponding page is not resident in memory and must therefore be retrieved from disk
    - **page fault handler** loads the appropriate page from disk and then returns

control to the instruction that caused the fault

## H6 Aborts

- Result from **unrecoverable fatal errors** (typically hardware errors)
- **Abort handlers NEVER** return control to the application program
  - returns control to an **abort** routine that **terminates the application program**

## H5 ¶8.1.3 Exceptions in Linux / x86-64 Systems

- There are ~256 different exception types
  - 0~31 --> defined by the Intel architects  $\iff$  identical for any x86-64 system
  - 32~255 --> **interrupts & traps** defined by the OS
- Examples of exceptions

## H6 Linux / x86-64 Faults and Aborts

- **Divide error** (exception 0)
  - occurs when an application attempts to divide by zero or when the result of a divide instruction is too big for the destination operand
  - Unix does not attempt to recover from divide errors --> **Abort**
  - Linux shells report as " **Floating exceptions** "
- **General protection fault** (exception 13)
  - when a program references an undefined area of virtual memory or because the program attempts to write to a read-only text segment
  - Linux does not attempt to recover from this fault
  - Linux shells report as " **Segmentation faults** "
- **Page fault** (exception 14)
  - handler maps the appropriate page of virtual memory on disk into a page of physical memory, then **restarts the faulting instruction**
- **Machine check** (exception 18)
  - when fatal hardware error is detected during the execution of the faulting instruction
  - **Never** return control to the application program

## H6 Linux / x86-64 System Calls

- **Linux** provides **hundreds of system calls** that application programs use when they **want to request services from the kernel**
    - reading a file, writing a file, creating a new process
    - each system call has a **unique integer number** that corresponds to an **offset in a jump table** in the kernel (not the same as the exception table)
  - C programs can invoke any system call **directly** by using the **syscall** function
  - In general, it is better (easier) to use **wrapper functions (System-level functions)**
    - package up the arguments, trap to the kernel with the appropriate system call instruction, pass the return status of the system call back to the calling program
    - all arguments to Linux system calls are **passed through general-purpose registers**. By convention,
      - %rax - syscall number
      - %rdi, %rsi, %rdx, %r10, %r8, %r9 -- arguments
- > On return from the system call, %rcx, %r11 are destroyed, and **%rax contains the return value**
- negative return value (-4096~-1) **indicates an error**

[Example C code]

```
1  int main() {
2      write(1, "hello, world\n", 13);
3      _exit(0);
4  }
```

- Arguments to **write** system-level function
  - 1st argument --> sends the output to **stdout**
  - 2nd argument --> sequence of bytes to write
  - 3rd argument --> number of bytes to write

[Implementing C code above directly with **Linux system calls**]

```
1  .section .data
```

```

2  string:
3      .ascii "hello, world\n"
4  string_end:
5      .equ len, string_end - string
6  .section .text
7  .globl main
8  main:
9      # First, call write(1, "hello, world\n", 13)
10     movq $1, %rax      # write is system call 1
11     movq $1, %rdi      # Arg1: stdout = descriptor 1
12     movq $string, %rsi # Arg2: hello world string
13     movq $len, %rdx     # Arg3: string length
14     syscall            # Make the system call
15
16     # Next, call _exit(0)
17     movq $60, %rax     # _exit is system call 60
18     movq $0, %rdi      # Arg1: exit status is 0
19     syscall            # Make the system call

```

### H3 §8.2 Processes

- When running a program on a modern system, we are presented with the illusion that our program is the **only one currently running** in the system --> provided by **Process**
  - our program seems to have exclusive use of the processor / memory
  - processor appears to execute the instructions without interruption
  - the code & data appear to be the only objects in the system's memory
- Process - an instance of a program in execution
  - each program in the system runs in the context of some process
    - context consists of the state that the program needs to run correctly
      - program's code & data stored in memory + stack + contents of its general purpose registers + program counter + environment variables + set of open file descriptors
- Each time a user **runs a program**, the **shell creates a new process** and then **runs the executable object file** in the context of this new process



- Key abstractions provided by a process
  - **Independent logical control flow** -- providing the illusion that our program has **exclusive use of the processor**
  - **Private address space** -- providing the illusion that our program has **exclusive use of the memory space**

## H5 ¶8.2.1 Logical Control Flow

- Illusion that the program has exclusive use of the processor (even though many other programs are typically running concurrently on the system)
- **However**, when we use a debugger to **single-step the execution** of our program, we would observe a series of **program counter (PC)** values that **corresponded exclusively** to instructions contained in the program's executable object file or in shared objects linked into the program dynamically at run time

--> sequence of PC values is known as **logical control flow (logical flow)**

- Suppose there is a system that runs three processes, then the single physical control flow of the processor is partitioned into three logical flows (one for each process)
  - processes take turns using the processor
  - each process executes a portion of its flow and then is **preempted** (temporarily suspended) while other processes take their turns
    - can be observed by measuring the elapsed time of each instruction

## H5 ¶8.2.2 Concurrent Flows

- **Concurrent flow** -- a logical flow whose execution overlaps in time with another flow (two flows are said to **run concurrently** )
- **Multitasking** -- notion of a process taking turns with other processes (sometimes referred to as **time slicing** )
- **Time slice** -- each time period that a process executes a portion of its flow
- If two flows **overlap in time** --> they are **concurrent**  
(independent of the number of processor cores or computers that the flows are running on)
- If two flows are **running concurrently on different processor cores or computers** --> **parallel flows** (running in parallel, or have parallel execution)

## H5 ¶8.2.3 Private Address Space

- Process provides each program with the **illusion** that it has exclusive use of the **system's address space**
  - provides with its own **private address space**
    - **cannot** in general be **read or written by any other process**
- Each space **has the same general organization**
  - **Bottom portion** of the address space is reserved for the **user program**
    - usual code, data, heap, stack segments
    - code segment always begins at address 0x400000
  - **Top portion** of the address space is reserved for the **kernel** (memory-resident part of the OS)
    - code, data, stack
    - being used when the **kernel executes instructions** on behalf of the process (system calls)

## H5 ¶8.2.4 User and Kernel Modes

- **Mode bit** in some control register -- characterizes the **privileges** that the process currently enjoys
  - Mode bit **set** -- process running in **kernel mode** (**supervisor mode**)
    - **can execute any instruction** in the instruction set and **access any memory location** in the system
  - Mode bit **not set** -- process running in **user mode**
    - **not allowed to** execute **privileged instructions**
      - halt the processor, change the mode bit, initiate an I/O operation,
      - not allowed to directly reference code or data in the kernel area of the address space
    - --> **fatal protection fault**
    - --> User programs **MUST** access kernel code and data **indirectly** via the **system call interface**

- Application code is **initially in user mode**
  - To change from **user mode** to **kernel mode** is by exceptions
    - **interrupt, fault, trapping system call**
      - When  $\exists$  exception, **Control passes to the exception handler**, the processor changes the mode to **kernel mode**
        - Handler runs in kernel mode and when it **returns to the application code**, the processor changes the mode **back to user mode**
- **/proc** filesystem
  - allows **user mode processes to access the contents of kernel data structures**
  - exports the contents of many kernel data structures as a hierarchy of text files that can be read by user programs
- **/sys** filesystem
  - additional low-level information about system buses and devices

## H5 ¶8.2.5 Context Switches

- Context switch -- **multitasking** using a higher-level form of exceptional control flow
- **Kernel** maintains a context for each process
  - state that the **kernel needs to restart** a preempted process
    - general-purpose registers + floating-point registers + program counter + user's stack + status registers + kernel's stack + kernel data structures
      - **page table** (characterizes the address space), **process table** (contains information about the current process), **file table** (contains information about the files that the process has opened)
- At certain points during the execution of a process, the **kernel can decide to preempt the current process and restart a previously preempted process** ( scheduling ) -- handled by code in the kernel ( scheduler )
- Scheduled -- when the kernel selects a new process to run
  - > **Kernel preempts the current process and transfers control to the new process** using a mechanism called **Context switch**

1. saves the context of the current process
2. restores the saved context of some previously preempted process
3. passes control to this newly restored process

--> can occur while the **kernel is executing a system call** on behalf of the user

- When the system call blocks, because it is **waiting** for some event to occur, then the **kernel can put the current process to sleep and switch to another process**
  - e.g. **read** system call requires a disk access,  
kernel can opt to perform a context switch and **run another process instead of waiting for the data** to arrive from the disk
  - e.g. **sleep** system call  
explicit request to put the calling process to sleep  
(even if a system call does not block, the kernel can decide to perform a context switch rather than return control to the calling process)
- **Context switch** can also occur as a result of an **interrupt**
  - e.g. when the disk sends an interrupt to signal that data have been transferred from disk to memory (return back to original process (prior to **read** )

### H3 §8.3 System Call Error Handling

- When Unix system-level functions encounter an error, they typically **return -1** and set global integer variabel **errno** to indicate what went wrong
- Programmers should **always check** for errors

[Example C code]

```
1  if ((pid = fork()) < 0) {
2      fprintf(stderr, "fork error: %s\n", strerror(errno));
3      exit(0);
4  }
```

--> **strerror** function returns a text string that describes the error associated with a particular value of **errno**

--> can simplify this code by defining the following **error-reporting function**:

```

1 void unix_error(char *msg) /* Unix-style error */
2 {
3     fprintf(stderr, "%s: %s\n", msg, strerror(errno));
4     exit(0);
5 }
6
7 /* In some function */
8 if ((pid = fork()) < 0)
9     unixerror("fork error");

```

- We can simplify our code even further by using **error-handling wrappers**
  - For a given base function **foo**, we define a **wrapper** function **Foo** with identical arguments but with the **first letter** of the name **capitalized**
    - wrapper function..
      - calls the base function
      - checks for errors
      - terminates if there are any problems

[Example C code]

```

1 pid_t Fork(void) {
2     pid_t pid;
3
4     if ((pid = fork()) < 0)
5         unix_error("Fork error");
6     return pid;
7 }
8
9 /* In some function */
10 pid = Fork()

```

### H3 §8.4 Process Control

#### H5 ¶8.4.1 Obtaining Process IDs

- Each process has a **unique positive (nonzero) process ID (PID)**.
- **getpid** function returns the **PID** of the calling process
- **getppid** function returns the **PID** of its **parent**

```

1 #include <sys/types.h>
2 #include <unistd.h>
3
4 pid_t getpid(void);
5 pid_t getppid(void);

```

--> return an integer value of type **pid\_t** (defined in types.h as an **int**)

## H5 ¶8.4.2 Creating and Terminating Processes

- Think of a **process** as being in one of **three** states
  1. **Running**
    - process is either **executing on the CPU** or **waiting** to be executed and will eventually be **scheduled by the kernel**

### 2. **Stopped**

- execution of the process is **suspended** and **will not be scheduled**
- as a result of receiving a
  - **SIGSTOP, SIGTSTP, SIGTTIN, SIGTTOU** signal
- stopped until it receives a
  - **SIGCONT** signal, at which point it becomes running again

--> signal - a form of software interrupt

### 3. **Terminated**

- process is **stopped permanently**
- three reasons to be terminated
  1. receiving a signal whose default action is to terminate the process
  2. returning from the main routine
  3. calling the **exit** function

```

1 #include <stdlib.h>
2 void exit(int status); // Function does not return

```

--> terminates the process with an **exit status** of status

(other way is to return an integer value from the main routine)

- **Parent process** creates a new running **child process** by calling the **fork** function

```

1  #include <sys/types.h>
2  #include <unistd.h>
3
4  pid_t fork(void); // Returns 0 to child, PID to parent, -1 on
    error

```

- Newly created process is almost, but not quite, **identical to the parent**
- The child gets an **identical (but separate)** copy of the parent's
  - **User-level virtual address space**
    - code
    - data segments
    - heap
    - shared libraries
    - user stack
  - **Open file descriptors**
    - child can read and write any files that were open in the parent when it called **fork**
- **Major difference** between the parent and the newly created child
  - **Different PIDs**
- **Fork** function is quite confusing since it is **called once**, but it **returns twice**
  - once in the **calling process**
    - returns the **PID of the child**
  - once in the newly created **child process**
    - returns a value of **0**

--> provides an unambiguous way to tell whether the program is executing in the parent or the child

[Example C code]

```

1  int main() {
2      pid_t pid;
3      int x = 1;
4

```

```

5  pid = Fork();
6  if (pid == 0) {      /* Child */
7      printf("child : x = %d\n", ++x);
8      exit(0);
9  }
10
11  /* Parent */
12  printf("parent: x = %d\n", --x);
13  exit(0);
14  }

```

--> will return

```

1  linux> ./fork
2  parent : x = 0
3  child : x = 2

```

- Aspects of **fork** functions

1. **Call once, return twice**

- called once by the parent, but it returns twice
  - once to the parent and once to the newly created child

2. **Concurrent execution**

- parent and child are **separate processes** that run **concurrently**
- parent process completes its statement first, then followed by the child (with example above)
  - on another system, the reverse might be true

3. **Duplicate but separate address spaces**

- when we stop both the parent and the child immediately after the **fork** function returned in each process, the address space of each process is **identical**
  - each process has the same user stack, local variable values, heap, global variable values, and code
- HOWEVER**, since the parent and the child are **separate processes**,
  - they each have their **own private address spaces**



- any subsequent changes that a parent or child makes to **x** are private and are not reflected in the memory of other process

#### 4. Shared files

- when the parent calls **fork**, the **stdout** file is open and directed to the screen
- the child inherits this file, and its output is also directed to the screen

- **Nested fork calls**

```

1  int main(){
2      Fork();
3      Fork();
4      printf("hello\n");
5      exit(0);
6  }

```

--> prints hello four times

### H5 ¶8.4.3 Reaping Child Processes

- When a process terminates for any reason, the kernel does not remove it from the system immediately
  - Instead, the process is **kept** around in a terminated state **until it is reaped** by its parent
    - When the parent **reaps** the terminated process, at which point it ceases to exist
    - a terminated process that has not yet been reaped is called a **zombie**
- When a parent process terminates, the kernel arranges for the **init** process to become the adopted parent of any **orphaned children**
  - **init** process has a PID of 1
    - is created by the kernel during system start-up, **never terminates**, and is the **ancestor of every process**
  - If a parent process terminates without reaping its zombie children, then the kernel arranges for the **init** process to reap them
- Long-running programs (shells, servers) should **always reap their zombie children** ( ∴

even though zombies are not running, they still consume system memory resources)

- A process waits for its children to terminate or stop by calling the **waitpid** function

```
1 #include <sys/types.h>
2 #include <sys/wait.h>
3
4 pid_t waitpid(pid_t pid, int *statusp, int options);
5 // Returns PID of child if OK, 0 (if WNOHANG), or -1 on error
```

- By default (options = 0), **waitpid** suspends execution of the calling process until a child process in its **wait set** terminates
  - If a process in the wait set has **already terminated at the time of the call**, then **waitpid** returns immediately
- > In either case, **waitpid** returns the PID of the terminated child that caused **waitpid** to return
- > at this point, the terminated child has been reaped and the kernel removes all traces of it from the system

## H6 Determining the Members of the Wait Set

- Members of the wait set are determined by the **pid** argument
  - If **pid** > 0, then the wait set is the **singleton child process** whose process ID is equal to **pid**
  - If **pid** = -1, then the wait set consists of all of the parent's child processes

## H6 Modifying the Default Behavior

- Default behavior can be modified by setting **options** to various combinations of the **WNOHANG**, **WUNTRACED**, **WCONTINUED** constants
- **WNOHANG**
  - **return immediately** (with a return value of **0**) if **none** of the child processes in the wait set has **terminated yet**
  - useful when you **want to continue doing useful work while waiting** for a child to

terminate

- **WUNTRACED**
  - suspend execution of the calling process until a process in the wait set becomes **either terminated or stopped**
  - return the **PID** of the **terminated or stopped child** that caused the return (by default, only for the terminated child)
  - useful when you want to **check for both terminated and stopped children**
- **WCONTINUED**
  - suspend execution of the calling process until a running process in the wait set is **terminated** or until a stopped process in the wait set **has been resumed** by the receipt of a **SIGCONT** signal
- **WNOHANG | WUNTRACED**
  - return **immediately** , with a return value of **0** , if **none** of the children in the wait set **has stopped or terminated** , or with a return value to the **PID** of **one of the stopped or terminated children**

## H6 Checking the Exit Status of a Reaped Child

- If the **statusp** argument is non-NULL, the **waitpid** encodes **status information about the child** that caused the return in **status** (value pointed to by **statusp** )
- **wait.h** include file defines several macros for interpreting the **status** argument
- **WIFEXITED** (status) -- returns true if the child terminated normally, via a call to **exit** or a return
- **WEXITSTATUS** (status) -- returns the exit status of a normally terminated child (is only defined if **WIFEXITED** () returned true)
- **WIFSIGNALED** (status) -- returns true if the child process terminated because of a signal that was not caught
- **WTERMSIG** (status) -- returns the number of the signal that caused the child process to terminate (is only defined if **WIFSIGNALED** () returned true)
- **WIFSTOPPED** (status) -- returns true if the child that caused the return is currently stopped
- **WSTOPSIG** (status) -- returns the number of the signal that caused the child to stop (is

only defined if **WIFSTOPPED()** returned true)

- **WIFCONTINUED** (status) -- returns true if the child process was restarted by receipt of a **SIGCONT** signal

## H6 Error Conditions

- If the calling process has no children, then **waitpid** returns -1 and sets **errno** to **ECHILD**
- If the **waitpid** function was interrupted by a signal, then it returns -1 and sets **errno** to **EINTR**

## H6 The wait Function

- **wait** function is a simpler version of **waitpid**

```
1  #include <sys/types.h>
2  #include <sys/wait.h>
3
4  pid_t wait(int *statusp);
```

--> calling **wait(&status)** is equivalent to calling **waitpid(-1, &status, 0)**

## H6 Examples of Using waitpid

[Example C code]

```
1  #include "csapp.h"
2  #define N 2
3
4  int main() {
5      int status, i;
6      pid_t pid;
7
8      /* Parent creates N children */
9      for (i = 0; i < N; i++)
10         if ((pid = Fork()) == 0) /* Child */
11             exit(100+i);
12
13     /* Parent reaps N children in no particular order */
14     while ((pid = waitpid(-1, &status, 0)) > 0) {
```

```

15     if (WIFEXITED(status))
16         printf("child %d terminated normally with exit status=%d\n",
pid, WEXITSTATUS(status));
17     else
18         printf("child %d terminated abnormally\n", pid);
19 }
20
21 /* The only normal termination is if there are no more children */
22 if (errno != ECHILD)
23     unix_error("waitpid error");
24
25 exit(0);
26 }

```

- Uses **waitpid** to wait for all of its  $N$  children to terminate
- In L10, the parent creates each of the  $N$  children, and in L11, each child exits with a unique exit status
- In L14, the parent waits for all of its children to terminate by using **waitpid** as the test condition of a **while** loop
  - since the first arg = -1, the call to **waitpid** blocks until an arbitrary child has terminated
- L15 checks the exit status of the child
  - if the child terminated normally (by calling the **exit** function), then the parent extracts the exit status and prints it on **stdout**
- When all of children have been reaped, the next call to **waitpid** returns **-1**, and sets **errno** to **ECHILD**
  - L22 checks that the **waitpid** function terminated normally, and prints an error message otherwise

--> this program reaps its children in no particular order

=> **nondeterministic** behavior (makes reasoning about concurrency so difficult)

--> **never** assume that one outcome will always occur

--> each possible outcome is equally likely

- **However, simple change** can eliminate this nondeterminism in the output order

```

1  #include "csapp.h"
2  #define N 2
3

```

```

4  int main() {
5      int status, i;
6      pid_t pid[N], retpid;
7
8      /* Parent creates N children */
9      for (i = 0; i < N; i++)
10         if ((pid[i] = Fork()) == 0)    /* Child */
11             exit(100+i);
12
13     /* Parent reaps N children in order */
14     i = 0;
15     while ((retpid = waitpid(pid[i++], &status, 0)) > 0) {
16         if (WIFEXITED(status))
17             printf("child %d terminated normally with exit
status=%d\n", retpid, WEXITSTATUS(status));
18         else
19             printf("child %d terminated abnormally\n", retpid);
20     }
21
22     /* The only normal termination is if there are no more
children */
23     if (errno != ECHILD)
24         unix_error("waitpid error");
25
26     exit(0);
27 }

```

--> In L10, the parent stores the PID's of its children in order and then waits for each child in this same order by calling **waitpid** with the appropriate **PID** in the first argument (rather than -1 for arbitrary)

## H5 ¶8.4.4. Putting Processes to Sleep

- **sleep** function suspends a process for a specified period of time

```

1  #include <unistd.h>
2  unsigned int sleep(unsigned int secs);
3                                     // Returns seconds left to sleep

```

- returns 0 -- if the requested amount of time has elapsed
- returns number of seconds still left to sleep

- if **sleep** function returns prematurely because it was interrupted by a **signal**
- **pause** function puts the calling function to sleep until a signal is received by the process

```
1 #include <unistd.h>
2 int pause(void);           // Always returns -1
```

## H5 ¶8.4.5 Loading and Running Programs

- **execve** function **loads and runs** a new program in the context of the current process

```
1 #include <unistd.h>
2 int execve(const char *filename, const char *argv[], const char
  *envp[]);
3           // Does not return if OK; returns -1 on error
```

- loads and runs the executable object file **filename** with the argument list **argv** and the environment variable list **envp**
  - **returns** to the calling program **only if** there is an **error** (not being able to find **filename**)
    - called once but never returns
- **argv** -- points to a null-terminated array of pointers, each of which points to an **argument string** (e.g. ls, -lt, ...)
- **argv[0]** -- name of the executable object file
- **envp** -- points to a null-terminated array of pointers to **environment variable strings**, each of which is a name-value pair of the form (e.g. PWD=/usr/---, USER = ---)( **name = value** )
- After **execve** loads **filename**, it calls the **start-up code**
  - sets up the stack
  - passes control to the main routine of the new program which has a prototype of the form

```

1  int main(int argc, char *argv[], char *envp[]);
2
3  /* Which is equivalent to */
4  int main(int argc char **argv, char **envp);

```

- Stack organization when **main** begins executing (from bottom to top)
  - argument & environment strings
  - null-terminated array of pointers, each of which points to an environment variable string
    - global variable **environ** points to the first of these pointers, **envp[0]**
  - null-terminated **argv[]**, with each element pointing to an argument string on the stack
  - [TOP] - stack frame for the system start-up function **libc\_start\_main**
- Three arguments are stored in registers
  1. **argc** (in %rdi) -- gives the number of non-null pointers in the **argv[]** array
  2. **argv** (in %rsi) -- points to the first entry in the **argv[]** array
  3. **envp** (in %rdx) -- points to the first entry in the **envp[]** array
- Several functions for manipulating the environment array

```

1  #include <stdlib.h>
2  char *getenv(const char *name);
3      // Returns: pointer to name if it exists, NULL if no match

```

- **getenv** function searches the environment array for a string **name=value**
  - If found, it returns a pointer to **value**
  - otherwise, returns **NULL**



```

1 #include <stdlib.h>
2 int setenv(const char *name, const char *newvalue, int
   overwrite);
3 // Returns: 0 on success, -1 on error
4
5 void unsetenv(const char *name);
6 // Returns: nothing

```

- **setenv** replaces **oldvalue** with **newvalue** only if **overwrite** is nonzero
  - If **name** does not exist, then **setenv** adds **name=newvalue** to the array
- **unsetenv** deletes **name=oldvalue** if it exists

### H3 §8.5 Signals

- **Linux signal** -- allows processes and the kernel to interrupt other processes
  - small message that **notifies a process** that an event of some type has occurred in the system
  - each signal type corresponds to some kind of system event
    - **low-level hardware exceptions** are processed by the **kernel's exception handlers** and **would not normally be visible to user processes**
      - > **Signals expose** the occurrence of such exceptions to user processes

e.g. if a process attempts to divide by zero, then the kernel sends it a SIGFPE signal

- other signals correspond to higher-level software events in the kernel or in other user processes

e.g. typing Ctrl+C while a process is running in the foreground, then the kernel sends a SIGINT (no. 2) to each process in the foreground process group

### H5 ¶8.5.1 Signal Terminology

- Transfer of a signal to a destination process occurs in two distinct steps

#### 1. Sending a signal

- kernel **sends** (delivers) a signal to a destination process by **updating some state** in the context of the destination process when...
  1. the kernel has **detected a system event** such as a divide-by-zero error or the termination of a child process

2. a process has invoked the **kill** function to explicitly request the kernel to send a signal to the destination process

## 2. Receiving a signal

- destination process **receives** a signal **when it is forced by the kernel to react** in some way to the delivery of the signal
- the process **can** either **ignore** the signal, **terminate**, or **catch** the signal by executing a user-level function called a **signal handler**
- **Pending signal** -- a signal that has been sent but not yet received
  - at any point in time, there can be **at most one** pending signal of a particular type
  - if a process has a **pending signal of type  $k$** , then any **subsequent** signals of type  $k$  sent to that process are **not queued**  
--> simply discarded
    - a process can selectively **block** the receipt of certain signals
      - when a signal is blocked, it can still be delivered, but the resulting **pending signal will not be received until the process unblocks the signal**
  - is received at most once
    - for each process, the kernel maintains the set of pending signals in the **pending bit vector** & set of blocked signals in the **blocked bit vector**

## H5 ¶8.5.2 Sending Signals

## H6 Process Groups

- Every process belongs to exactly one **process group**, identified by a positive integer **process group ID**
- **getpgrp** function returns the process group ID of the current process

```
1  #include <unistd.h>
2  pid_t getpgrp(void);
3  // Returns: process group ID of calling process
```

- By default, a child process belongs to the same process group as its parent. A process can change the process group of itself or another process by using the **setpgid** function

```
1  #include <unistd.h>
2  int setpgid(pid_t, pid, pid_t pgid);
3  // Returns: 0 on success, -1 on error
```

--> changes the process group of process **pid** to **pgid**

- If **pid** is zero, the PID of the current process is used
- If **pgid** is zero, the PID of the process specified by **pid** is used for the process group ID

## H6 Sending Signals with the /bin/kill Program

- **/bin/kill** program sends an arbitrary signal to another process

```
1  linux> /bin/kill -9 15213
```

sends signal 9 (**SIGKILL**) to process 15213

- A **negative PID** causes the signal to be sent to every process in process group PID

```
1  linux> /bin/kill -9 -15213
```

sends a SIGKILL signal to every process in process group 15213

## H6 Sending Signals from the Keyboard

- Unix shells use the abstraction of a **job** to represent the processes that are created as a result of evaluating a single command line
- There is at most one foreground job and zero or more background jobs
- For example,

```
1  linux> ls | sort
```

- creates a foreground job consisting of two processes connected by a **Unix pipe**
  - one running the **ls** program
  - other running the **sort** program

--> the shell creates a separate process group for each job

--> typically, the process group ID is taken from one of the parent processes in the job

- Typing Ctrl+C at the keyboard causes the kernel to send a **SIGINT** signal to every process in the foreground process group
  - by default, the result is to terminate the foreground job
- Similarly, Ctrl+Z causes the kernel to send a **SIGTSTP** signal to every process in the foreground process group
  - by default, the result is to stop the foreground job

## H6 Sending Signals with the kill Function

- Processes send signals to other processes (including themselves) by calling the **kill** function

```
1 #include <sys/types.h>
2 #include <signal.h>
3 int kill(pid_t pid, int sig);
4           // Returns: 0 if OK, -1 on error
```

- If **pid** > 0, then the **kill** function sends signal number **sig** to process **pid**
- If **pid** = 0, then **kill** sends signal **sig** to every process in the group of the calling process, including the calling process itself
- If **pid** < 0, then **kill** sends signal **sig** to every process in process group **|pid|**

[Example C code] -- example of a parent that uses the **kill** function to send a **SIGKILL** signal to its child

```
1 #include "csapp.h"
2 int main() {
3     pid_t pid;
4
5     /* Child sleeps until SIGKILL signal received, then dies */
6     if ((pid = Fork()) == 0) {
7         Pause();
8         printf("control should never reach here!\n");
```

```

 9      exit(0);
10  }
11
12  /* Parent sends a SIGKILL signal to a child */
13  kill(pid, SIGKILL);
14  exit(0);
15  }

```

## H6 Sending Signals with the alarm Function

- Process can send **SIGALRM** signal to itself by calling the **alarm** function

```

1  #include <unistd.h>
2  unsigned int alarm(unsigned int secs);
3  // Returns: remaining seconds of previous alarm, or 0 if no
   previous alarm

```

- arranges for the kernel to send a **SIGALRM** signal to the calling process in **secs** seconds
  - if **secs** = 0, then no new alarm is scheduled
  - in any event, the call to **alarm** cancels any pending alarms and returns the number of seconds remaining until any pending alarm was due to be delivered, or 0 if there were no pending alarms

## H5 ¶8.5.3 Receiving Signals

- When the kernel switches a process  $p$  from kernel mode to user mode (returning from a system call or completing a context switch), the kernel checks the set of **unblocked pending signals** for  $p$ 
  - If the set is **empty** (usual case)
    - **kernel** passes control to the next instruction ( $I_{next}$ ) in the logical control flow of  $p$
  - If the set is **nonempty**
    - **kernel** chooses some signal  $k$  in the set (typically the smallest) and forces  $p$  to **receive** signal  $k$ 
      - --> triggers some **action** by the process
      - once the process completes the action, then control passes back to the next instruction in the logical control flow of  $p$

- **Default action**
  - process terminates
  - process terminates and dumps core
  - process stops (suspends) until restarted by a **SIGCONT** signal
  - process ignores the signal
- A process **can modify the default action** associated with a signal by using the **signal** function
  - exceptions -- **SIGSTOP, SIGKILL** (default actions cannot be changed)

```

1  #include <signal.h>
2  typedef void (*sighandler_t)(int);
3
4  sighandler_t signal(int signum, sighandler_t handler);
5  // Returns: pointer to previous handler if OK, SIG_ERR on
   error (does not set errno)

```

- **signal** function can change the action associated with a signal **signum** in one of three ways:
  - if **handler** is **SIG\_IGN**, then signals of type **signum** are ignored
  - if **handler** is **SIG\_DFL**, then the action for signals of type **signum** reverts to the default action
  - Otherwise, **handler** is the address of a user-defined function, called a **signal handler**, that will be called whenever the process receives a signal of type **signum**
    - changing the default action by passing the address of a handler to the **signal** function is known as **installing the handler**
    - invocation of the handler is called **catching the signal**
    - execution of the handler is referred to as **handling the signal**
- When a process catches a signal of type  $k$ , the handler installed for signal  $k$  is invoked with a single integer argument set to  $k$ 
  - allows the same handler function to catch different types of signals
- When the handler executes its **return** statement, control (usually) passes back to the instruction in the control flow where the process was interrupted by the receipt of the signal

[Example C code] -- program that catches the SIGINT signal that is sent whenever the user types Ctrl+C at the keyboard

```
1  #include "csapp.h"
2
3  void sigint_handler(int sig) /* SIGINT handler */
4  {
5      printf("Caught SIGINT!\n");
6      exit(0);
7  }
8
9  int main() {
10     /* Install the SIGINT handler */
11     if (signal(SIGINT, sigint_handler) == SIG_ERR)
12         unix_error("signal error");
13
14     pause(); /* Wait for the receipt of a signal */
15
16     return 0;
17 }
```

- default action for **SIGINT** is to immediately terminate the process
  - however, in this case, we modify the default behavior to catch the signal, print a message, and then terminate the process
- Signal handlers can be interrupted by other handlers

## H5 ¶8.5.4 Blocking and Unblocking Signals

- Implicit blocking mechanism
  - by default, the kernel blocks any pending signals of the type currently being processed by a handler
- Explicit blocking mechanism
  - applications can **explicitly** block and unblock selected signals using the **sigprocmask** function and its helpers

```

1  #include <signal.h>
2  int sigprocmask(int how, const sigset_t *set, sigset_t *oldset);
3  int sigemptyset(sigset_t *set);
4  int sigfillset(sigset_t *set);
5  int sigaddset(sigset_t *set, int signum);
6  int sigdelset(sigset_t *set, int signum);
7      // Returns: 0 if OK, -1 on error
8  int sigismember(const sigset_t *set, int signum);
9      // Returns: 1 if member, 0 if not, -1 on error

```

- **sigprocmask** changes the set of currently blocked signals (the **blocked** bit vector)
  - specific behavior depends on the value of **how**
    - **SIG\_BLOCK** -- add the signals in **set** to **blocked** ( **blocked = blocked | set** )
    - **SIG\_UNBLOCK** -- remove the signals in **set** from **blocked** ( **blocked = blocked & ~set** )
    - **SIG\_SETMASK** -- **blocked = set**
  - If **oldset** is non-NULL, the previous value of the **blocked** bit vector is stored in **oldset**
- **sigemptyset** -- initializes **set** to the empty set
- **sigfillset** -- adds every signal to **set**
- **sigaddset** -- adds **signum** to set
- **sigdelset** -- deletes **signum** from set
- **sigismember** -- returns 1 if **signum** is a member of set, and 0 if not

[Example C code]

```

1  sigset_t mask, prev_mask;
2
3  Sigemptyset(&mask);
4  Sigaddset(&mask, SIGINT);
5
6  /* Block SIGINT and save previous blocked set. */
7  Sigprocmask(SIG_BLOCK, &mask, &prev_mask);
8
9      // Code region that will not be interrupted by SIGINT
10 /* Restore previous blocked set, unblocking SIGINT */
11 Sigprocmask(SIG_SETMASK, &prev_mask, NULL);

```



## H5 ¶8.5.5 Writing Signal Handlers

- What makes signal handling difficult?
  1. Handlers run concurrently with the main program and share the same global variables
    - > can interfere with the main program and with other handlers
  2. how/when signals are received is often counterintuitive
  3. different systems can have different signal-handling semantics

## H6 Safe Signal Handling

- If a handler and the main program access the same global data structure concurrently, then the results can be unpredictable and often fatal
- To avoid concurrency errors,
  0. Keep handlers as simple as possible
    - make handler set a global flag and return immediately
    - let all processing performed by the main program
  1. Call only async-signal-safe functions in your handlers
    - **async-signal-safe** functions can be **safely called** from a **signal handler** either because
      - it is **reentrant** (e.g. accesses only local vars)
      - it cannot be interrupted by a signal handler
    - The only safe way to generate output from a signal handler is to use **write** function
      - **printf, sprintf** are not safe
    - **\_exit** is an async-signal-safe variant of **exit**
  2. Save and restore **errno**
    - many of the Linux async-signal-safe functions set **errno** when they return with an error
    - calling such functions inside a handler might interfere with other parts of the program that rely on **errno**
      - save **errno** to a local variable on entry to the handler and restore it before the handler returns

- only necessary if the handler returns
- not necessary if the handler terminates the process by calling **\_exit**

### 3. Protect accesses to shared global data structures by blocking all signals

- if a handler shares a global data structure with the main program or with other handlers, handlers and main program should **temporarily block all signals** while accessing that data structure

### 4. Declare global variables with **volatile**

- tell the compiler **not to cache a variable** by declaring it with the **volatile** type qualifier

e.g. `volatile int g;`

--> forces the compiler to read the value of *g* from memory each time it is referenced in the code

### 5. Declare **flags** with **sig\_atomic\_t**

- handler records the receipt of the signal by writing to a global **flag**
- C provides an integer data type, **sig\_atomic\_t** -- reads and writes are guaranteed to be **atomic (uninterruptible)** because they can be implemented with a single instruction

`volatile sig_atomic_t flag;`

- can safely read from and write to **sig\_atomic\_t** variables without temporarily blocking signals

## H6 Correct Signal Handling

- Pending signals are not queued
- **pending** bit vector contains exactly one bit for each type of signal

--> there can be **at most one** pending signal of any particular type

- if two signals of type *k* are sent to a destination process while signal *k* is blocked (because the destination process is currently executing a handler for signal *k*), then the second signal is **discarded**

[Example C code]

```

1 void handler1(int sig) {
2     int olderrno = errno;
3
4     if ((waitpid(-1, NULL, 0)) < 0)
5         sio_error("waitpid error");
6     Sio_puts("Handler reaped child\n");
7     Sleep(1);
8     errno = olderrno;
9 }
10
11 int main() {
12     int i, m;
13     char buf[MAXBUF];
14
15     if (signal(SIGCHLD, handler1) == SIG_ERR)
16         unix_error("signal error");
17
18     /* Parent creates children */
19     for (i = 0; i < 3; i++) {
20         if (Fork() == 0) {
21             printf("Hello from child %d\n", (int)getpid());
22             exit(0);
23         }
24     }
25
26     /* Parent waits for terminal input and then processes it */
27     if ((n = read(STDIN_FILENO, buf, sizeof(buf))) < 0)
28         unix_error("read");
29
30     printf("Parent processing input\n");
31     while(1)
32         ///
33
34     exit(0);
35 }

```

- This code failed to account for the fact that signals are not queued
- First signal is received and caught by the parent
  - Second signal is delivered while the handler is still processing the first signal
    - > second signal added to the set of pending signals (not received yet since SIGCHLD signals are blocked by the SIGCHLD handler)
  - Third signal arrives while the handler is still processing the first signal

--> Third signal **discarded**

--> Signals cannot be used to count the occurrence of events in other processes

[Example C code w/o forementioned error]

```
1 void handler2(int sig) {
2     int olderrno = errno;
3
4     while (waitpid(-1, NULL, 0) > 0) {
5         Sio_puts("Handler reaped child\n");
6     }
7     if (errno != ECHILD)
8         Sio_error("waitpid error");
9     sleep(1);
10    errno = olderrno;
11 }
```

--> When all of the children have been reaped, the next call to waitpid returns -1, and sets errno to ECHILD

--> L7 checks that the waitpid function terminated normally

## H6 Portable Signal Handling

- Different systems have different signal-handling semantics
  - Semantics of the **signal** function varies
    - some older Unix systems restore the action for signal  $k$  to its default after signal  $k$  has been caught by a handler
      - handler must explicitly reinstall itself each time it runs
  - System calls can be interrupted
    - **slow system calls** -- system calls that can potentially block the process for a long period
      - on some older versions of Unix, slow system calls that are interrupted when a handler catches a signal do not resume when the signal handler returns
        - > return immediately to the user with an error condition and **errno** set to **EINTR**
      - programmers must include code that manually restarts interrupted system calls

- To deal with abovementioned issues, the Posix standard defines the **sigaction** function, which allows users to **clearly specify the signal-handling semantics** they want when they install a handler

```
1 #include <signal.h>
2 int sigaction(int signum, struct sigaction *act, struct sigaction
  *oldact);
3 // Returns: 0 if OK, -1 on error
```

--> inconvenient since it requires the user to set the entries of a complicated structure

--> define a wrapper function: **Signal** that calls **sigaction**

```
1 handler_t *Signal (int signum, handler_t *handler) {
2     struct sigaction action, old_action;
3
4     action.sa_handler = handler;
5     sigemptyset(&action.sa_mask);
6     action.sa_flags = SA_RESTART;
7
8     if (sigaction(signum, &action, &old_action) < 0)
9         unix_error("Signal error");
10    return (old_action.sa_handler);
11 }
```

- **Signal** wrapper installs a signal handler with the following signal handling semantics:
  - only signals of the type currently being processed by the handler are blocked
  - as with all signal implementations, signals are not queued
  - interrupted system calls are automatically restarted whenever possible
  - once the signal handler is installed, it remains installed until **Signal** is called with a **handler** argument of either **SIG\_IGN** or **SIG\_DFL**